

EE 5561: Image Processing and Applications

Lecture 9

Mehmet Akçakaya

Recap of Last Lecture

- Linear algebra review
 - Wavelet transforms via linear algebra
 - Haar wavelets, unitary
 - Energy compaction
 - Another transform with good energy compaction (DCT)
 - Random variables & vectors review
- Today
 - Start on image restoration

Enhancement vs. Restoration


- Enhancement
 - Pleasing images
 - No model for degradation
 - Ad hoc procedures
 - Visual metric (mostly)
- Restoration
 - “Inverting” a degradation
 - Model exists
 - Quantitative evaluation

Enhancement vs. Restoration


- Degradations are common in imaging systems
 - Either by design or physical limitations
- Restoration relies on optimization procedures
 - We already saw least squares as a basic example
- Related to math field “inverse problems”
- Sometimes also called “computational imaging”

Image Denoising

- First problem we will consider in this module



- Goal: Recover \mathbf{x} from \mathbf{y}
- Our first attempt will be based on Wiener filter
 - Original version (with random processes) by Wiener in 1930s
 - Easy to implement a linear shift invariant filter for recovery



The original formulation considers signals of possibly infinite extent (not vectors)

Image Denoising – Wiener Filter

- In matrix form we will want

$$\hat{\mathbf{x}} = \mathbf{G}\mathbf{y} \quad \text{(just a linear system)}$$

- Wiener's criterion: Minimize the average mean squared error
 - \mathbf{x} itself is an instance of a random vector \mathbf{X}
 - I am going to drop the notation for random vectors, and use them interchangeably, since the context is clear here
 - i.e. minimize $\mathbb{E}[\|\mathbf{x} - \hat{\mathbf{x}}\|_2^2]$

Image Denoising – Wiener Filter

– We have

$$\begin{aligned}\epsilon^2 &= \mathbb{E}[\|\mathbf{x} - \hat{\mathbf{x}}\|_2^2] = \mathbb{E}[(\mathbf{x} - \hat{\mathbf{x}})^H (\mathbf{x} - \hat{\mathbf{x}})] \\ &= \mathbb{E}\left[\text{tr}((\mathbf{x} - \hat{\mathbf{x}})^H (\mathbf{x} - \hat{\mathbf{x}}))\right] \\ &= \mathbb{E}\left[\text{tr}((\mathbf{x} - \hat{\mathbf{x}})(\mathbf{x} - \hat{\mathbf{x}})^H)\right] \\ &= \text{tr}\left(\mathbb{E}[(\mathbf{x} - \hat{\mathbf{x}})(\mathbf{x} - \hat{\mathbf{x}})^H]\right) \\ &= \text{tr}\left(\mathbb{E}[\mathbf{x}\mathbf{x}^H - \hat{\mathbf{x}}\mathbf{x}^H - \mathbf{x}\hat{\mathbf{x}}^H + \hat{\mathbf{x}}\hat{\mathbf{x}}^H]\right)\end{aligned}$$

$\text{tr}(\mathbf{A})$: sum of diagonal elements of square matrix \mathbf{A}

since $\text{tr}(\mathbf{AB}) = \text{tr}(\mathbf{BA})$

since $\text{tr}(\cdot)$ and $\mathbb{E}(\cdot)$ are linear

– We need other assumptions

- \mathbf{x} and \mathbf{n} are uncorrelated (important)
- \mathbf{n} is zero-mean (can be worked around but for easier notation)

Image Denoising – Wiener Filter

– Then

$$\epsilon^2 = \text{tr} \left(\mathbb{E} [\mathbf{x} \mathbf{x}^H - \hat{\mathbf{x}} \mathbf{x}^H - \mathbf{x} \hat{\mathbf{x}}^H + \hat{\mathbf{x}} \hat{\mathbf{x}}^H] \right)$$

▪ Here $\mathbb{E} [\mathbf{x} \mathbf{x}^H] = \mathbf{R}_x$

$$\begin{aligned} \mathbb{E} [\hat{\mathbf{x}} \mathbf{x}^H] &= \mathbb{E} [(\mathbf{G} \mathbf{y}) \mathbf{x}^H] = \mathbb{E} [\mathbf{G} (\mathbf{x} + \mathbf{n}) \mathbf{x}^H] \\ &= \mathbb{E} [\mathbf{G} \mathbf{x} \mathbf{x}^H + \mathbf{G} \mathbf{n} \mathbf{x}^H] \\ &= \mathbf{G} \left(\mathbb{E} [\mathbf{x} \mathbf{x}^H + \mathbf{n} \mathbf{x}^H] \right) \\ &= \mathbf{G} \left(\mathbb{E} [\mathbf{x} \mathbf{x}^H] + \mathbb{E} [\mathbf{n}] \mathbb{E} [\mathbf{x}^H] \right) \\ &= \mathbf{G} \left(\mathbb{E} [\mathbf{x} \mathbf{x}^H] \right) = \mathbf{G} \mathbf{R}_x \end{aligned}$$

Image Denoising – Wiener Filter

– Then

$$\epsilon^2 = \text{tr} \left(\mathbb{E} [\mathbf{x}\mathbf{x}^H - \hat{\mathbf{x}}\mathbf{x}^H - \mathbf{x}\hat{\mathbf{x}}^H + \hat{\mathbf{x}}\hat{\mathbf{x}}^H] \right)$$

▪ Here $\mathbb{E} [\mathbf{x}\mathbf{x}^H] = \mathbf{R}_x$

$$\mathbb{E} [\hat{\mathbf{x}}\mathbf{x}^H] = \mathbf{G}\mathbf{R}_x$$

$$\begin{aligned} \mathbb{E} [\hat{\mathbf{x}}\hat{\mathbf{x}}^H] &= \mathbb{E} [\mathbf{G}\mathbf{y}\mathbf{y}^H\mathbf{G}^H] = \mathbf{G}\mathbb{E} [\mathbf{y}\mathbf{y}^H]\mathbf{G}^H \\ &= \mathbf{G}\mathbb{E} [(\mathbf{x} + \mathbf{n})(\mathbf{x} + \mathbf{n})^H]\mathbf{G}^H \\ &= \mathbf{G}\mathbb{E} [\mathbf{x}\mathbf{x}^H + \mathbf{n}\mathbf{x}^H + \mathbf{x}\mathbf{n}^H + \mathbf{n}\mathbf{n}^H]\mathbf{G}^H \\ &= \mathbf{G}(\mathbf{R}_x + \mathbf{R}_n)\mathbf{G}^H \end{aligned}$$

▪ So we have

$$\epsilon^2 = \text{tr} \left(\mathbf{R}_x - \mathbf{G}\mathbf{R}_x - \mathbf{R}_x^H\mathbf{G}^H + \mathbf{G}\mathbf{R}_x\mathbf{G}^H + \mathbf{G}\mathbf{R}_n\mathbf{G}^H \right)$$

Image Denoising – Wiener Filter

- We want to minimize this error with respect to \mathbf{G}

$$\epsilon^2 = \text{tr}(\mathbf{R}_x - \mathbf{G}\mathbf{R}_x - \mathbf{R}_x^H \mathbf{G}^H + \mathbf{G}\mathbf{R}_x \mathbf{G}^H + \mathbf{G}\mathbf{R}_n \mathbf{G}^H)$$

- i.e. we want to set $\frac{\partial \epsilon^2}{\partial \mathbf{G}} = 0$

- How?

$$\frac{\partial \epsilon^2}{\partial \mathbf{G}} = \begin{bmatrix} \frac{\partial \epsilon^2}{\partial g_{11}} & \frac{\partial \epsilon^2}{\partial g_{12}} & \cdots \\ \vdots & & \end{bmatrix}$$

- This becomes trickier since the matrices can be complex

Image Denoising – Wiener Filter

- Complex matrix calculus (IEEE TSP, 55(6): 2740-46)

$$\frac{\partial}{\partial \mathbf{G}} \text{tr}(\mathbf{A}\mathbf{G}) = \mathbf{A}^T$$

$$\frac{\partial}{\partial \text{conj}(\mathbf{G})} \text{tr}(\mathbf{A}\mathbf{G}) = \mathbf{0}$$

$$\frac{\partial}{\partial \mathbf{G}} \text{tr}(\mathbf{G}^H \mathbf{A}) = \mathbf{0}$$

$$\frac{\partial}{\partial \text{conj}(\mathbf{G})} \text{tr}(\mathbf{G}^H \mathbf{A}) = \mathbf{A}$$

$$\frac{\partial}{\partial \mathbf{G}} \text{tr}(\mathbf{G}\mathbf{A}\mathbf{G}^H) = \mathbf{G}^H \mathbf{A}$$

$$\frac{\partial}{\partial \text{conj}(\mathbf{G})} \text{tr}(\mathbf{G}\mathbf{A}\mathbf{G}^H) = \mathbf{G}\mathbf{A}$$

- Based on these easier to set

$$0 = \frac{\partial \epsilon^2}{\partial \text{conj}(\mathbf{G})}$$

$$= \frac{\partial}{\partial \text{conj}(\mathbf{G})} \text{tr}(\mathbf{R}_x) - \frac{\partial}{\partial \text{conj}(\mathbf{G})} \text{tr}(\mathbf{G}\mathbf{R}_x) - \frac{\partial}{\partial \text{conj}(\mathbf{G})} \text{tr}(\mathbf{R}_x^H \mathbf{G}^H) + \frac{\partial}{\partial \text{conj}(\mathbf{G})} \text{tr}(\mathbf{G}(\mathbf{R}_x + \mathbf{R}_n)\mathbf{G}^H)$$

$$= \mathbf{0} - \mathbf{0} - \mathbf{R}_x^H + \mathbf{G}(\mathbf{R}_x + \mathbf{R}_n)$$

$$= -\mathbf{R}_x + \mathbf{G}(\mathbf{R}_x + \mathbf{R}_n)$$

Image Denoising – Wiener Filter

- We want to minimize this error with respect to \mathbf{G}

$$\epsilon^2 = \text{tr} \left(\mathbf{R}_x - \mathbf{G}\mathbf{R}_x - \mathbf{R}_x^H \mathbf{G}^H + \mathbf{G}\mathbf{R}_x \mathbf{G}^H + \mathbf{G}\mathbf{R}_n \mathbf{G}^H \right)$$

- Leads to $\mathbf{G} = \mathbf{R}_x (\mathbf{R}_x + \mathbf{R}_n)^{-1}$
- Assuming white noise gives $\mathbf{R}_n = \sigma^2 \mathbf{I}$
- \mathbf{R}_x is Hermitian positive semi-definite (it is a correlation matrix) \rightarrow diagonalize it

$$\mathbf{R}_x = \mathbf{U} \mathbf{\Lambda} \mathbf{U}^H \quad \text{with } \mathbf{U} = [\mathbf{u}_1 \dots \mathbf{u}_n] \text{ unitary}$$

and $\Lambda_{ii} = \lambda_i$ diagonal

Image Denoising – Wiener Filter

- We want to minimize this error with respect to \mathbf{G}

$$\epsilon^2 = \text{tr} \left(\mathbf{R}_x - \mathbf{G}\mathbf{R}_x - \mathbf{R}_x^H \mathbf{G}^H + \mathbf{G}\mathbf{R}_x \mathbf{G}^H + \mathbf{G}\mathbf{R}_n \mathbf{G}^H \right)$$

- With these

$$\begin{aligned} \mathbf{G} &= \mathbf{R}_x (\mathbf{R}_x + \mathbf{R}_n)^{-1} \\ &= \mathbf{U} \mathbf{\Lambda} \mathbf{U}^H (\mathbf{U} \mathbf{\Lambda} \mathbf{U}^H + \sigma^2 \mathbf{I})^{-1} \\ &= \mathbf{U} \mathbf{\Lambda} \mathbf{U}^H \left(\mathbf{U} (\mathbf{\Lambda} + \sigma^2 \mathbf{I}) \mathbf{U}^H \right)^{-1} \\ &= \mathbf{U} \mathbf{\Lambda} (\mathbf{\Lambda} + \sigma^2 \mathbf{I})^{-1} \mathbf{U}^H \\ &= \sum_{k=1}^n \frac{\lambda_k}{\lambda_k + \sigma^2} \mathbf{u}_k \mathbf{u}_k^H \end{aligned}$$

Image Denoising – Wiener Filter

– How does this work in reality?

- \mathbf{R}_x is generally not available (or equivalently \mathbf{U} and $\mathbf{\Lambda}$ are not available)
- Need to estimate it
- One approach¹ is to approximate \mathbf{U} with a pre-determined transform (e.g. wavelet), then estimate $\mathbf{\Lambda}$ from transform coefficients

$$\lambda_i = \begin{cases} |\theta(i)|^2 & \text{for } i = 1, \dots, n_s \\ 0 & \text{otherwise} \end{cases} \quad (\text{based on some threshold for energy preservation})$$

- Alternative: For large n , \mathbf{R}_x diagonalized by IDFT (\sim random process limit)
- This ties in with how for wide-sense stationary random processes, spectral density is the Fourier transform of correlation function

Image Denoising – Wiener Filter

– How does this work in reality?

- Or patch processing
- We already mentioned that block processing is popular in image processing (and we'll see more)
 - Define a small neighborhood
 - Assume similar signal content here
 - In this case, this means: image pixels in small neighborhood come from the same distribution
 - Calculate sample mean, standard deviation
 - For pixels $\{y_1, \dots, y_n\}$

$$\hat{\mu}_y = \frac{1}{n} \sum_{k=1}^n y_k$$

$$\hat{\sigma}_y = \sqrt{\frac{\sum_{k=1}^n (y_k - \hat{\mu}_y)^2}{n - 1}}$$

Bessel's correction

Performance Measures

- We saw our first restoration tool, designed to minimize a certain “loss” based on a statistical model
 - We modeled both noise and image as instances of random vectors
 - Used an expected l_2 loss
- We will see such losses or performance measures for the remainder of the course
 - Here we'll briefly review some common ones

Performance Measures

- Mean squared error (MSE):

$$\text{MSE} = \frac{1}{N} ||\mathbf{x} - \hat{\mathbf{x}}||_2^2 \quad \mathbf{x} \in \mathbb{R}^N \text{ (vectorized image)}$$

or

$$\text{MSE} = \mathbb{E} [||\mathbf{x} - \hat{\mathbf{x}}||_2^2] \quad \text{if } \mathbf{x} \text{ is a random vector}$$

- This is per-pixel average error
- Its units: square of the original image unit
- We sometimes want the same units as the image, especially in quantitative imaging modalities (e.g. CT)

- Root mean squared error $\text{RMSE} = \sqrt{\text{MSE}}$

Performance Measures

- Normalized MSE (NMSE)

- MSE and RMSE both depend on how \mathbf{x} is scaled (e.g. is it between 0-1 or 0-255?)
- NMSE tackles this

$$\text{NMSE} = \frac{\text{MSE}}{\frac{1}{N} \|\mathbf{x}\|_2^2} = \frac{\|\mathbf{x} - \hat{\mathbf{x}}\|_2^2}{\|\mathbf{x}\|_2^2}$$

or statistical version

$$\text{NMSE} = \frac{\text{MSE}}{\mathbb{E}[\|\mathbf{x}\|_2^2]}$$

- $\text{SNR} = \frac{1}{\text{NMSE}} \quad \text{or} \quad 10 \log_{10} \frac{1}{\text{NMSE}}$

Performance Measures

- Peak SNR (PSNR)

$$\text{PSNR} = 10 \log_{10} \frac{\max_k |x_k|^2}{\text{MSE}}$$

- Very common for evaluating image quality

- Worst case (l_∞)

$$\max_k |\hat{x}_k - x_k| = \|\mathbf{x} - \hat{\mathbf{x}}\|_\infty$$

- l_1 error

$$\frac{1}{N} \|\mathbf{x} - \hat{\mathbf{x}}\|_1 = \frac{1}{N} \sum_{k=1}^N |\hat{x}_k - x_k|$$

Performance Measures

– Structural Similarity Index (SSIM)

- Main idea: Previous metrics work on quantifying point-wise differences between images. But our perception system identifies structural information & differences in those.
- Need a metric to capture this.
- Relies on 3 features: luminance, contrast, structure
- Luminance: Averaging over all pixels. For image \mathbf{x} , $\mu_x = \frac{1}{N} \sum_{k=1}^N x_k$
- Contrast: Standard deviation across pixels. $\sigma_x = \left(\frac{1}{N-1} \sum_{k=1}^N (x_k - \mu_x)^2 \right)^{\frac{1}{2}}$
- Structure: Normalized input signal $(\mathbf{x} - \mu_x) / \sigma_x$
- Then we build functions to compare each of these between two given images...
- And combine these information
- After some work, the standard formula is
$$\text{SSIM}(\mathbf{x}, \mathbf{y}) = \frac{(2\mu_x\mu_y + C_1)(2\sigma_{xy} + C_2)}{(\mu_x^2 + \mu_y^2 + C_1)(\sigma_x^2 + \sigma_y^2 + C_2)}$$

Performance Measures

- Structural Similarity Index (SSIM)
 - This is a global metric
 - It is better to apply the metrics regionally (on patches), and then average these
 - Technically called mean structural similarity index
 - Leads to weighted averages (e.g. with a Gaussian weighting) instead of simple averaging for luminance, contrast and structure
 - Final result is the average across all patches

Performance Measures

- Not covered now
 - Perceptual losses
- These and SSIM will be important later in the ML/AI module too