

EE 5561: Image Processing and Applications


Lecture 4

Mehmet Akçakaya

Last Lecture Recap

- Fourier analysis review
 - Fourier series for periodic (or finite-extent) images
 - Fourier transform as a more general tool
 - How Fourier analysis describes images
 - Low-frequency content → contrast information, most of the energy of the image
 - High-frequency content → edges and other sharper structures
- Sampling & interpolation
 - The tools that allow us to move to discrete-space
- 2D discrete-space signals introduction

2D Discrete-Space Systems



- Recall Kronecker delta
$$\delta[m, n] = \begin{cases} 1 & \text{if } m = n = 0 \\ 0 & \text{otherwise} \end{cases}$$
- Properties analogous to Dirac delta for continuous-space signals


$$\sum_{m,n} \delta[m, n] = 1$$

$$\delta[m - m_0, n - n_0] \cdot f[m, n] = f[m_0, n_0] \cdot \delta[m - m_0, n - n_0] \quad \text{sampling}$$

$$\sum_{m,n} (\delta[m - m_0, n - n_0] \cdot f[m, n]) = f[m_0, n_0] \quad \text{sifting}$$

$$\delta[m, n] = \delta_x[m] \cdot \delta_y[n] \quad \text{separable}$$

2D Discrete-Space Systems



- Example: moving average over a 3×3 window


$$\begin{aligned} g[m, n] &= \frac{1}{9} \sum_{k=m-1}^{m+1} \sum_{l=n-1}^{n+1} f[k, l] \\ &= \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[m-k, n-l] \end{aligned}$$

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

- Example: segmentation based on a simple threshold

$$g[m, n] = \begin{cases} 1 & |f[m, n]| > T \\ 0 & \text{otherwise} \end{cases}$$

2D Discrete-Space Systems



- Example: downsampling (trivial compression)

$$g[m, n] = f[2m, 2n]$$

- Example: upsampling

$$g[m, n] = \begin{cases} f[\frac{m}{2}, \frac{n}{2}] & m, n \text{ even} \\ 0 & \text{otherwise} \end{cases}$$

not ideal!

should interpolate...

2D Discrete-Space Systems

- Properties:
 - Amplitude-based: linearity, stability, invertibility
 - Spatial-based: causality, separability, shift-invariance,...

- Similar to continuous-space cases

- Linearity

$$\mathcal{S}(\alpha f_1 + \beta f_2) = \alpha \mathcal{S}(f_1) + \beta \mathcal{S}(f_2)$$

- e.g. thresholding (earlier) $f_1 + f_2 > T$ but $f_1 < T$ and $f_2 < T \rightarrow$ non-linear

- Stability: BIBO

- What about? $g[m, n] = g[m - 1, n - 1] + f[m, n]$
 - Take $f[m, n] = 1$
 - Bounded input does not produce bounded output

2D Discrete-Space Systems

- Invertibility
 - Downsampling? No
 - Upsampling? Yes
- Shift-invariance

$$\begin{aligned} \text{If } f[m, n] \rightarrow \mathcal{S} \rightarrow g[m, n] \quad \text{then} \\ f[m - m_0, n - n_0] \rightarrow \mathcal{S} \rightarrow g[m - m_0, n - n_0] \quad \forall m_0, n_0 \in \mathbb{Z}, \forall f \end{aligned}$$

- Moving average filter?

$$\begin{aligned} g[m, n] &= \mathcal{S}(f[m, n]) = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[m - k, n - l] \\ \mathcal{S}(f[m - m_0, n - n_0]) &= \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[m - m_0 - k, n - n_0 - l] \\ g[m - m_0, n - n_0] &= \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[m - m_0 - k, n - n_0 - l] \end{aligned}$$

← same!

2D Discrete-Space LSI Systems

- Same idea as in continuous space

- Instead using Kronecker delta as a building block

- For a general system: $\delta[m, n] \rightarrow \mathcal{S} \rightarrow h[m, n]$

- SI system: $\delta[m - k, n - l] \rightarrow \mathcal{S} \rightarrow h[m - k, n - l]$

- LSI system: $f[m, n] \rightarrow \mathcal{S} \rightarrow g[m, n]$

$$\begin{aligned} f[m, n] &= \sum_{k, l} f[k, l] \delta[m - k, n - l] \rightarrow \mathcal{S} \rightarrow \sum_{k, l} f[k, l] h[m - k, n - l] = g[m, n] \\ &= f[m, n] * h[m, n] \end{aligned}$$

discrete convolution

- This convolution operation will be fundamental
 - Will be used from linear filters to neural networks


Convolution Sum

- How to calculate the convolution sum? $\sum_{k,l} f[k,l]h[m-k,n-l]$
 - In signals & systems, view f and h as functions of $[k,l]$ for fixed $[m,n]$
 - Here $h[m-k,n-l]$ corresponds to a shifted & reversed version of $h[k,l]$
 - Process:
 - Fix $m,n \rightarrow$ Generate $h[m-k,n-l]$ with the appropriate shift & reversal
 - Multiply by $f[k,l]$ at each point \rightarrow sum over all k,l
 - Repeat over all m,n
 - Usually implemented as a sliding window
 - This is a little complicated... Why reverse for each m,n ? etc
- In practice, we just specify the “reversed” filter

Convolution Sum

– How to calculate the convolution sum?

▪ Pictorially



▪ An example:

3_0	3_1	2_2	1	0
0_2	0_2	1_0	3	1
3_0	1_1	2_2	2	3
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

Moving Average Filter Example

$f[m, n]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$g[m, n]$

Moving Average Filter Example

$f[m, n]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g[m, n]$

	0	10							

Moving Average Filter Example

$f[m, n]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g[m, n]$

	0	10	20						

Moving Average Filter Example

$f[m, n]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g[m, n]$

	0	10	20	30					

Moving Average Filter Example

$$f[m, n]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$g[m, n]$$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

Moving Average Filter Example

Original



3×3 kernel



5×5 kernel



Moving Average Filter Example

Original



Noisy




3×3 kernel






Convolution Sum

- Going back: What happens at the boundaries?
 - Already apparent from these examples: The output does not have the same size as the image by default
 - Why? Images have finite “support” on our computers
 - Support of image: $\{(m, n) \in \mathbb{Z}^2 : f[m, n] \neq 0\}$
 - With more edge processing




Convolution Sum

- What happens in practice at these boundaries?
 - Nothing is an option → gives the above result (usually not desirable)
 - Zero-padding



Convolution Sum

- What happens in practice at these boundaries?
 - Nothing is an option → gives the above result (usually not desirable)
 - Zero-padding
 - Mirror extension
 - 2D discrete cosine transform
 - Circular extension
 - 2D discrete Fourier transform



Convolution Sum

– Properties

$$f[m, n] * h[m, n] = h[m, n] * f[m, n]$$

commutative

$$(f[m, n] * g[m, n]) * h[m, n] = f[m, n] * (g[m, n] * h[m, n])$$

associative

$$f[m, n] * (h_1[m, n] + h_2[m, n]) = f[m, n] * h_1[m, n] + f[m, n] * h_2[m, n]$$

distributive

$$(f_1[m]f_2[n]) * (h_1[m]h_2[n]) = (f_1[m] *_{1D} h_1[m])(f_2[n] *_{1D} h_2[n])$$

separability

2D discrete-space Fourier Transform

- This is again motivated as eigenfunctions of LSI systems
 - Analogous to continuous-space
 - So we will skip to the derivation

$$F(\omega_x, \omega_y) = \sum_{m=-\infty}^{\infty} \sum_{n=-\infty}^{\infty} f[m, n] e^{-i(\omega_x m + \omega_y n)}$$

- Inverse is defined via an integral (we will skip this too)
- Properties
 - Periodic: $F(\omega_x, \omega_y) = F(\omega_x + 2\pi, \omega_y) = F(\omega_x, \omega_y + 2\pi)$
 - Transpose $f[m, n] \xleftrightarrow{\mathcal{F}} F(\omega_x, \omega_y)$
 $f[n, m] \xleftrightarrow{\mathcal{F}} F(\omega_y, \omega_x)$

2D discrete-space Fourier Transform

- Multiplication

$$f[m, n]g[m, n] \xleftrightarrow{\mathcal{F}} \frac{1}{(2\pi)^2} \int_{-\pi}^{\pi} \int_{-\pi}^{\pi} F(\lambda_x, \lambda_y) G(\omega_x - \lambda_x, \omega_y - \lambda_y) d\lambda_x d\lambda_y$$

2 π -periodic
convolution
in Fourier domain

- Convolution

$$f[m, n] * g[m, n] \xleftrightarrow{\mathcal{F}} F(\omega_x, \omega_y) G(\omega_x, \omega_y)$$

- Magnitude and phase

- In general $F(\omega_x, \omega_y)$ is complex, i.e. $F(\omega_x, \omega_y) = F_R(\omega_x, \omega_y) + iF_I(\omega_x, \omega_y)$
- As with other complex numbers, we can write this: $F(\omega_x, \omega_y) = |F(\omega_x, \omega_y)|e^{i\theta(\omega_x, \omega_y)}$
- Normally both available
- Some applications: One may be available (e.g. phase retrieval)

LSI Systems in Frequency Domain

- We have seen Fourier analysis over the past two lectures
 - So far, it helped us understand the frequency properties of images
 - But we can also use it for “system” or “filter” design
 - These will be simple pre-processing tools, e.g.
 - Noise reduction
 - Low-pass filter (why?)
 - Edge enhancement
 - High-pass filter (why?)
 - This understanding/intuition will come in handy when developing more complex image processing systems

LSI Systems in Frequency Domain


- Example (1D):

$$h[n] = \left[\underline{\frac{1}{4}} \quad \frac{1}{2} \quad \frac{1}{4} \right] \quad \text{underline denotes 0}^{\text{th}} \text{ location}$$

$$\begin{aligned} H(\omega) &= \sum_n h[n] e^{-i\omega n} = \frac{1}{4} + \frac{1}{2} e^{-i\omega} + \frac{1}{4} e^{-i2\omega} \\ &= \frac{1}{4} e^{-i\omega} (e^{i\omega} + 2 + e^{-i\omega}) = \frac{1}{2} e^{-i\omega} (1 + \cos \omega) \end{aligned}$$

$$|H(\omega)| = \frac{1}{2} (1 + \cos \omega)$$

$$\angle H(\omega) = -\omega$$




LSI Systems in Frequency Domain

- Example (1D):

$$h[n] = \begin{bmatrix} \frac{1}{4} & \frac{1}{2} & \frac{1}{4} \end{bmatrix}$$

$$|H(\omega)| = \frac{1}{2}(1 + \cos \omega)$$

$$\angle H(\omega) = -\omega$$



- What do these frequencies mean?

- Recall that DSFT is 2π periodic
 - 0 is still the lowest frequency
 - $\pm\pi$ are the high frequencies
 - This filter is 0 at $\pm\pi \rightarrow$ Low-pass filter

LSI Systems in Frequency Domain

- Example (Now in 2D):
$$h[m, n] = \begin{bmatrix} 0 & \frac{1}{4} & 0 \\ \frac{1}{4} & \frac{1}{2} & \frac{1}{4} \\ 0 & \frac{1}{4} & 0 \end{bmatrix}$$

$$\begin{aligned} H(\omega_x, \omega_y) &= \frac{1}{2} + \frac{1}{4}(e^{-i\omega_x} + e^{i\omega_x} + e^{-i\omega_y} + e^{i\omega_y}) \\ &= \frac{1}{2}(1 + \cos \omega_x + \cos \omega_y) \end{aligned}$$

- Check high-frequency performance

$$|H(\pi, 0)| = \frac{1}{2} \neq 0$$

- Not *ideal* low-pass

LSI Systems in Frequency Domain

- Example (Moving average):
$$h[m, n] = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$$H(\omega_x, \omega_y) = \frac{1}{9}(1 + 2 \cos \omega_x)(1 + 2 \cos \omega_y)$$

- Check high-frequency performance

$$|H(\pi, 0)| = \frac{1}{3} \neq 0$$

- Also not *ideal* low-pass

LSI Systems in Frequency Domain

- Example (Separable):
 - We already saw a better 1D low-pass filter (0 at $\pm\pi$)
 - We can build a 2D filter from it

$$h[m, n] = \begin{bmatrix} \frac{1}{4} \\ \frac{1}{2} \\ \frac{1}{4} \end{bmatrix} \begin{bmatrix} \frac{1}{4} & \frac{1}{2} & \frac{1}{4} \end{bmatrix} = \begin{bmatrix} \frac{1}{16} & \frac{1}{8} & \frac{1}{16} \\ \frac{1}{8} & \frac{1}{4} & \frac{1}{8} \\ \frac{1}{16} & \frac{1}{8} & \frac{1}{16} \end{bmatrix}$$

$$H(\omega_x, \omega_y) = \frac{1}{4}(1 + \cos \omega_x)(1 + \cos \omega_y)$$

- Check performance

$$H(0, 0) = 1$$

$$H(\pm\pi, \omega_y) = 0 = H(\omega_x, \pm\pi)$$

Yet Another Fourier Transform

- DSFT is useful for simple system analysis
- In reality: All our images have finite support (extent)
 - We do not need to sum from $\pm\infty$
- We need one final tool (which is what we will actually use)
 - Discrete Fourier transform (DFT)

Discrete-space Fourier Series

- To get to DFT, first consider a periodic discrete-space signal

$\tilde{x}[m, n]$ periodic with period (M, N)

- We follow the same process as in continuous-space, defining

$$\phi_{k,l}[m, n] = e^{i2\pi \left(k \frac{m}{M} + l \frac{n}{N} \right)}$$

- Do the projection argument...

$$\tilde{x}[m, n] = \sum_{k,l} c_{k,l} \phi_{k,l}[m, n]$$

Note $e^{-i2\pi(k+M)\frac{m}{M}} = e^{-i2\pi k \frac{m}{M}} \underbrace{e^{-i2\pi M \frac{m}{M}}}_{e^{-i2\pi m} \longleftarrow = 1 \text{ since } m \text{ is an integer}}$

- So we only need the summation w.r.t. (k,l) over one period!

Discrete-space Fourier Series

- We will not do the whole derivation, but this leads to the following:

$\tilde{x}[m, n]$ periodic with period (M, N)

$$\text{Define } \tilde{X}[k, l] = \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} \tilde{x}[m, n] e^{-i2\pi(k \frac{m}{M} + l \frac{n}{N})}$$


$$\text{Note } e^{-i2\pi(k+M)\frac{m}{M}} = e^{-i2\pi k \frac{m}{M}} \underbrace{e^{-i2\pi M \frac{m}{M}}}_{e^{-i2\pi m} \longleftarrow = 1 \text{ since } m \text{ is an integer}}$$

Thus $\tilde{X}[k, l]$ itself is periodic with period (M, N)

$$\text{and } \tilde{x}[m, n] = \frac{1}{MN} \sum_{k=0}^{M-1} \sum_{l=0}^{N-1} \tilde{X}[k, l] e^{i2\pi(k \frac{m}{M} + l \frac{n}{N})}$$

Discrete Fourier Transform

- What if $x[m,n]$ is not periodic, but finite $M \times N$ support (as most our images do)?
 - Make it periodic
 - (M,N) -point circular extension



Discrete Fourier Transform

- Circular extension

$$\tilde{x}[m, n] \triangleq x[m \bmod M, n \bmod N] \leftarrow \begin{array}{l} \text{depends on } x[m, n] \text{ for} \\ m \in \{0, \dots, M-1\} \\ n \in \{0, \dots, M-1\} \end{array}$$

- How to get the image back?

$$x[m, n] = \tilde{x}[m, n] R_{M \times N}[m, n] \leftarrow \text{rectangular window in image domain}$$



Discrete Fourier Transform

- Take the DSFT of the circular extension

$$\tilde{X}[k, l] = \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} \tilde{x}[m, n] e^{-i2\pi(k\frac{m}{M} + l\frac{n}{N})}$$

- DFT is when this is truncated

$$X[k, l] = \tilde{X}[k, l] R_{M \times N}[k, l]$$

- Everything we need to store is only defined over $M \times N$ regions

$$X[k, l] = \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} x[m, n] e^{-i2\pi(k\frac{m}{M} + l\frac{n}{N})} \quad k \in \{0, \dots, M-1\}, l \in \{0, \dots, N-1\}$$

$$x[m, n] = \frac{1}{MN} \sum_{k=0}^{M-1} \sum_{l=0}^{N-1} X[k, l] e^{i2\pi(k\frac{m}{M} + l\frac{n}{N})} \quad m \in \{0, \dots, M-1\}, n \in \{0, \dots, N-1\}$$

Discrete Fourier Transform


- Multiplication in DFT domain \rightarrow *Circular* convolution in image space

$$\begin{aligned} X[k, l]H[k, l] &\xleftarrow{DFT} x[m, n] \circledast h[m, n] \\ &= \sum_{m'=0}^{M-1} \sum_{n'=0}^{N-1} x[m', n'] h[(m - m') \bmod M, (n - n') \bmod N] \end{aligned}$$

- i.e. wrap-around
- Different than linear convolution (boundary conditions)
- If one wants to evaluate linear convolution (without wrap-around) \rightarrow need zero-padding (see earlier)
- Note unlike what you heard in EE 3015, in most image processing applications, it is faster to do convolutions in image space
 - The convolution filters usually have very small support (e.g. 3×3)

Discrete Fourier Transform

- Why DFT?
 - Allows us to analyze the spatial frequency spectrum of images
 - e.g. from last lecture: All generated with DFT
 - DFT can be implemented very fast
 - Using fast Fourier transform (FFT)



Brute force evaluation of DFT: $O((MN)^2)$ operations

FFT allows $O((MN) \log(\max(M, N)))$ operations from EE 4541


- FFT is what we use in practice

Fourier Magnitude and Phase

$f(x,y)$



$\text{angle}(F(u,v))$




$$\text{ifft2}(|G(u,v)|e^{i \cdot \text{angle}(F(u,v))})$$



$g(x,y)$



$\log|G(u,v)|$



Discrete Fourier Transform

- It will also be desirable to look at DFT in matrix-vector notation

- First let's do the 1D case
$$X[k] = \sum_{m=0}^{M-1} x[m] e^{-i2\pi(k \frac{m}{M})}$$

- Define matrix \mathbf{W}_M with entries $W_{k,n} = \left(e^{-i \frac{2\pi}{M}}\right)^{kn}$

- Then

$$\begin{bmatrix} X[0] \\ X[1] \\ \vdots \\ X[N-1] \end{bmatrix} = \begin{bmatrix} & \\ & \mathbf{W}_M \\ & \end{bmatrix} \begin{bmatrix} x[0] \\ x[1] \\ \vdots \\ x[N-1] \end{bmatrix}$$

- By properties of \mathbf{W}_M , we have

$$\mathbf{W}_M^H \mathbf{W}_M = M \mathbf{I}_M = \mathbf{W}_M \mathbf{W}_M^H \Rightarrow \mathbf{W}_M^{-1} = \frac{1}{M} \mathbf{W}_M^H$$

conjugate transpose

identity matrix

Discrete Fourier Transform

- Similarly easy to deduce a Parseval relation in this form

$$\begin{aligned}\mathbf{x}^H \mathbf{y} &= (\mathbf{W}_M^{-1} \mathbf{X})^H (\mathbf{W}_M^{-1} \mathbf{Y}) \\ &= \left(\frac{1}{M} \mathbf{W}_M^H \mathbf{X} \right)^H \left(\frac{1}{M} \mathbf{W}_M^H \mathbf{Y} \right) \\ &= \frac{1}{M^2} \mathbf{X}^H \underbrace{\mathbf{W}_M \mathbf{W}_M^H}_{M \mathbf{I}_M} \mathbf{Y} = \frac{1}{M} \mathbf{X}^H \mathbf{Y}\end{aligned}$$

Discrete Fourier Transform

- It will also be desirable to look at DFT in matrix-vector notation

- In 2D

$$\mathbf{x}_{\text{image}} = \begin{bmatrix} x[0, 0] & \cdots & x[0, N - 1] \\ \vdots & & \\ x[M - 1, 0] & \cdots & x[M - 1, N - 1] \end{bmatrix}$$

- By separability of the DFT, we have

$$\mathbf{X}_{\text{image}} = \mathbf{W}_M \mathbf{x}_{\text{image}} \mathbf{W}_N^T$$

Discrete Fourier Transform

- But often, it will be more convenient for us to *vectorize* the image

$$\mathbf{x}_{\text{vec}} = \text{vec}(\mathbf{x}_{\text{image}}) = \begin{bmatrix} x[0, 0] \\ \vdots \\ x[M-1, 0] \\ x[0, 1] \\ \vdots \\ x[M-1, N-1] \end{bmatrix}$$

- In that case

$$\mathbf{X}_{\text{vec}} = (\mathbf{W}_M \otimes \mathbf{W}_N) \mathbf{x}_{\text{vec}}$$

where

$$\mathbf{A} \otimes \mathbf{B} = \begin{bmatrix} a_{11}\mathbf{B} & \cdots & a_{1N}\mathbf{B} \\ \vdots & & \\ a_{m1}\mathbf{B} & \cdots & a_{mN}\mathbf{B} \end{bmatrix}$$

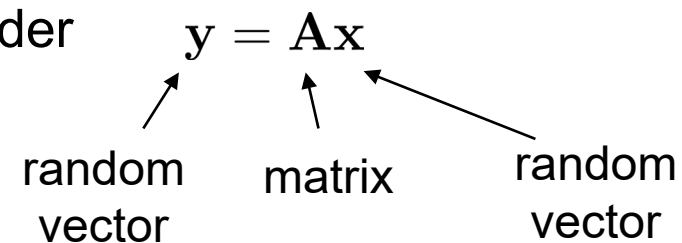
is the Kronecker product

← Note this is a linear operator
i.e. $\mathbf{X}_{\text{vec}} = \mathbf{L}\mathbf{x}_{\text{vec}}$
for linear operator
 $\mathbf{L} = \mathbf{W}_M \otimes \mathbf{W}_N$

Discrete Fourier Transform

- What happens to noise in the Fourier domain?

- Some basic probability

- Consider $y = Ax$


- Recall covariance matrix of \mathbf{y} is given as $\text{Cov}\{\mathbf{y}\} = \text{Cov}\{\mathbf{A}\mathbf{x}\} = \mathbf{A}\text{Cov}\{\mathbf{x}\}\mathbf{A}^H$

- Suppose random vector \mathbf{x} has uncorrelated entries with same variance, i.e. $\text{Cov}\{\mathbf{x}\} = \sigma^2 \mathbf{I}_N$

- Then its DFT $\mathbf{X} = \mathbf{W}_N \mathbf{x}$ has covariance matrix

$$\text{Cov}\{\mathbf{X}\} = \mathbf{W}_N \text{Cov}\{\mathbf{x}\} \mathbf{W}_N^H = \mathbf{W}_N \sigma^2 \mathbf{I}_N \mathbf{W}_N^H = \sigma^2 \mathbf{W}_N \mathbf{W}_N^H = \sigma^2 \cdot N \cdot \mathbf{I}_N$$

- Therefore uncorrelated in image domain \leftrightarrow uncorrelated in Fourier domain

- Note: This applies to any orthogonal/unitary transformation

Recap of 2D Signals & Systems

- Concept of signals & systems
 - Images are 2D signals
 - Image formation or processing systems
- Fourier analysis review
 - Continuous-space (Fourier series, Fourier transform)
 - Sampling & interpolation
 - Discrete-space (DSFT, DFT/FFT)
 - How Fourier analysis (via DFT) describes images
 - Low-frequency content → contrast information, most of the energy of the image
 - High-frequency content → edges and other sharper structures
- These intuitions & Fourier analysis will be useful for algorithm design
- Next: Start on image enhancement