

EE 5561: Image Processing and Applications

Lecture 17

Mehmet Akçakaya

Recap of Last Lecture

- Machine learning concepts

- Tasks, performance measures
- X : input data, Y : label
- f_{θ} : function we want to estimate/learn, parametrized by θ
- $\mathcal{L}(f_{\theta}(X), Y)$ loss function
- Sample mean loss
- Gradient descent
- Generalization gap
- Multi-layer networks
- Activation functions
- Neural networks

$$\arg \min_{\theta} \sum_{k=1}^n \mathcal{L}(f_{\theta}(x_k), y_k)$$

Minimizing the Loss

- Loss function is
$$\arg \min_{\boldsymbol{\theta}} \sum_{k=1}^n \mathcal{L}(f_{\boldsymbol{\theta}}(x_k), y_k)$$

- Optimization requires calculation of gradients with respect to $\boldsymbol{\theta}$
- Or more concretely, for our multi-layer feed-forward network

$$f_{\boldsymbol{\theta}}(\mathbf{x}) = \eta \left(\mathbf{W}^{[n]} \left(\dots \eta \left(\mathbf{W}^{[2]} \eta \left(\mathbf{W}^{[1]} \mathbf{x} + \mathbf{b}^{[1]} \right) + \mathbf{b}^{[2]} \right) \dots + \mathbf{b}^{[n]} \right) \right)$$

with respect to $\{\mathbf{W}^{[n]}, \dots, \mathbf{W}^{[1]}, \mathbf{b}^{[n]}, \dots, \mathbf{b}^{[1]}\}$

- How to calculate these derivatives?
 - Chain rule!

Minimizing the Loss

– Consider $y = f(g(h(x)))$

and let $u_1 = h(x), \quad u_2 = g(u_1), \quad y = f(u_2)$

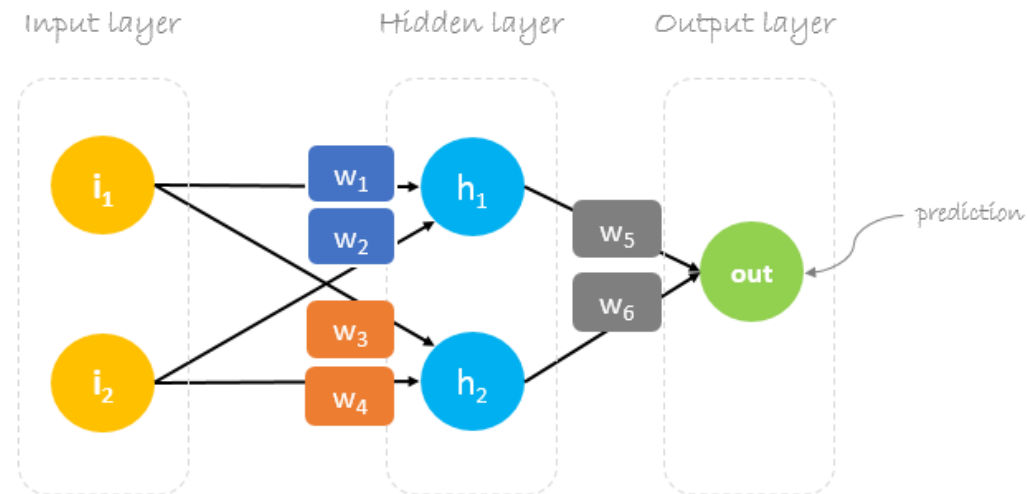
then
$$\frac{\partial f}{\partial x} \Big|_{x=c} = \frac{\partial f}{\partial u_2} \Big|_{u_2=g(u_1)} \frac{\partial g}{\partial u_1} \Big|_{u_1=h(c)} \frac{\partial h}{\partial x} \Big|_{x=c}$$

- In deep learning, chain rule is used in a procedure called backpropagation to calculate the gradient with respect to $\{\mathbf{W}^{[k]}, \mathbf{b}^{[k]}\}_k$ at each layer iteratively
- Basic idea

$$\frac{\partial f}{\partial x} \Big|_{x=c} = \frac{\partial f}{\partial u_2} \Big|_{u_2=g(h(c))} \frac{\partial g}{\partial u_1} \Big|_{u_1=h(c)} \frac{\partial h}{\partial x} \Big|_{x=c}$$

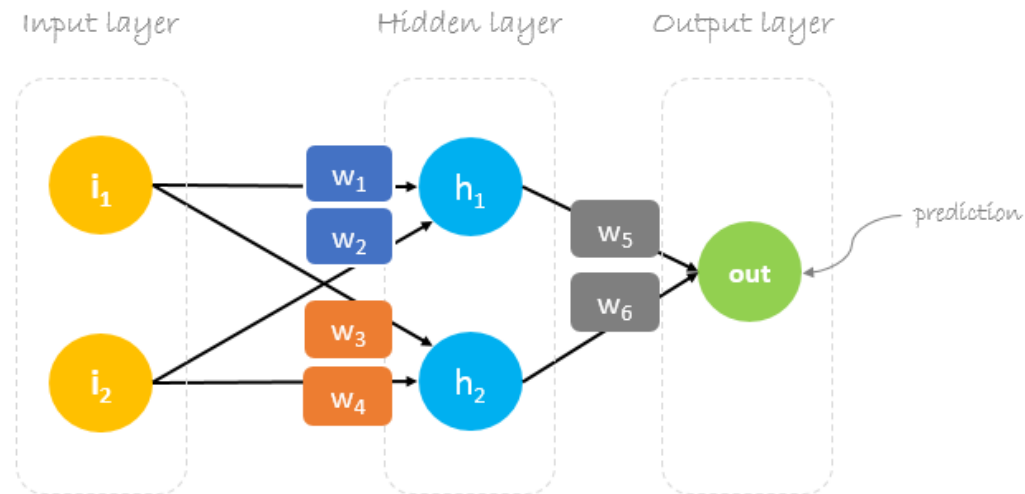
Backpropagation

Simple linear multi-layer network



Backpropagation

Simple linear multi-layer network

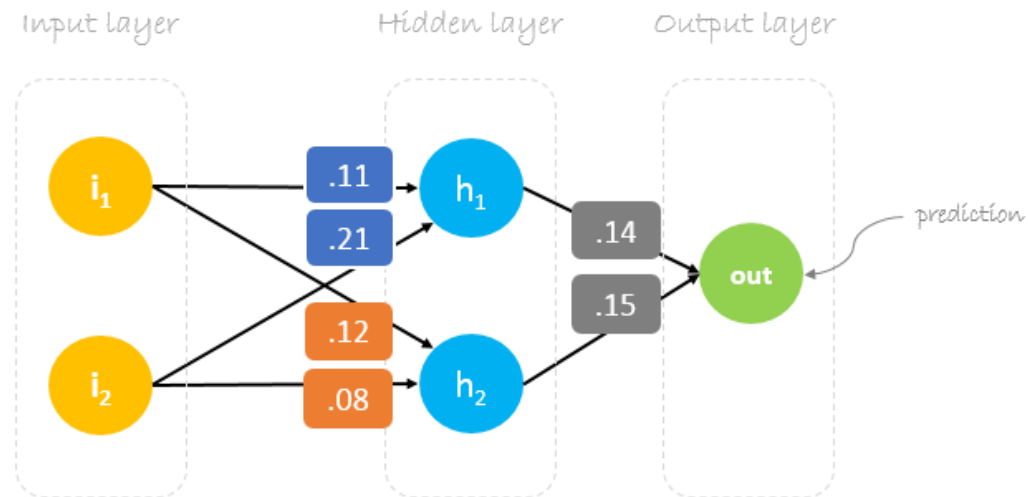


Assume initial weights

$w_1 = 0.11$, $w_2 = 0.21$, $w_3 = 0.12$, $w_4 = 0.08$, $w_5 = 0.14$ and $w_6 = 0.15$

Backpropagation

Simple linear multi-layer network



The updates depend on the database & loss function

Backpropagation

Let's use one sample in the dataset

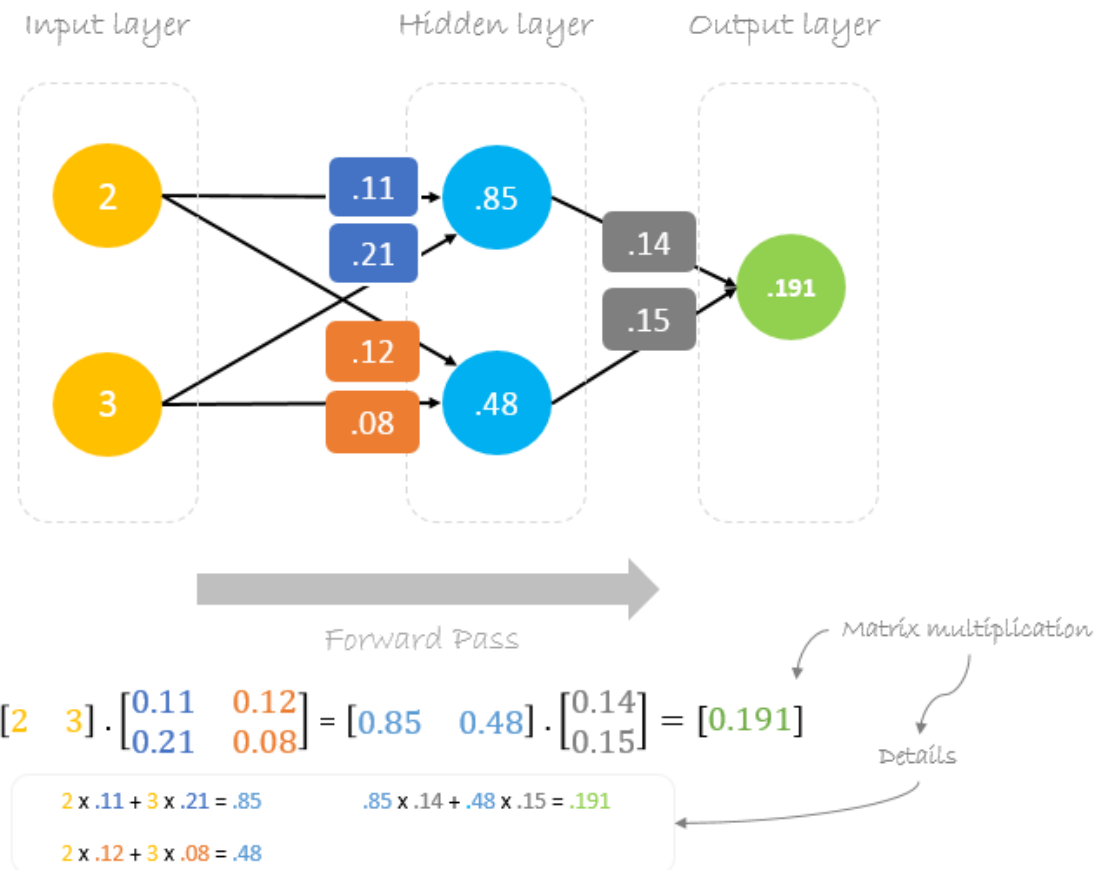


and use MSE loss

$$\text{MSE} = \frac{1}{2} (\text{predicted} - \text{label})^2 \triangleq (y - y^*)^2$$

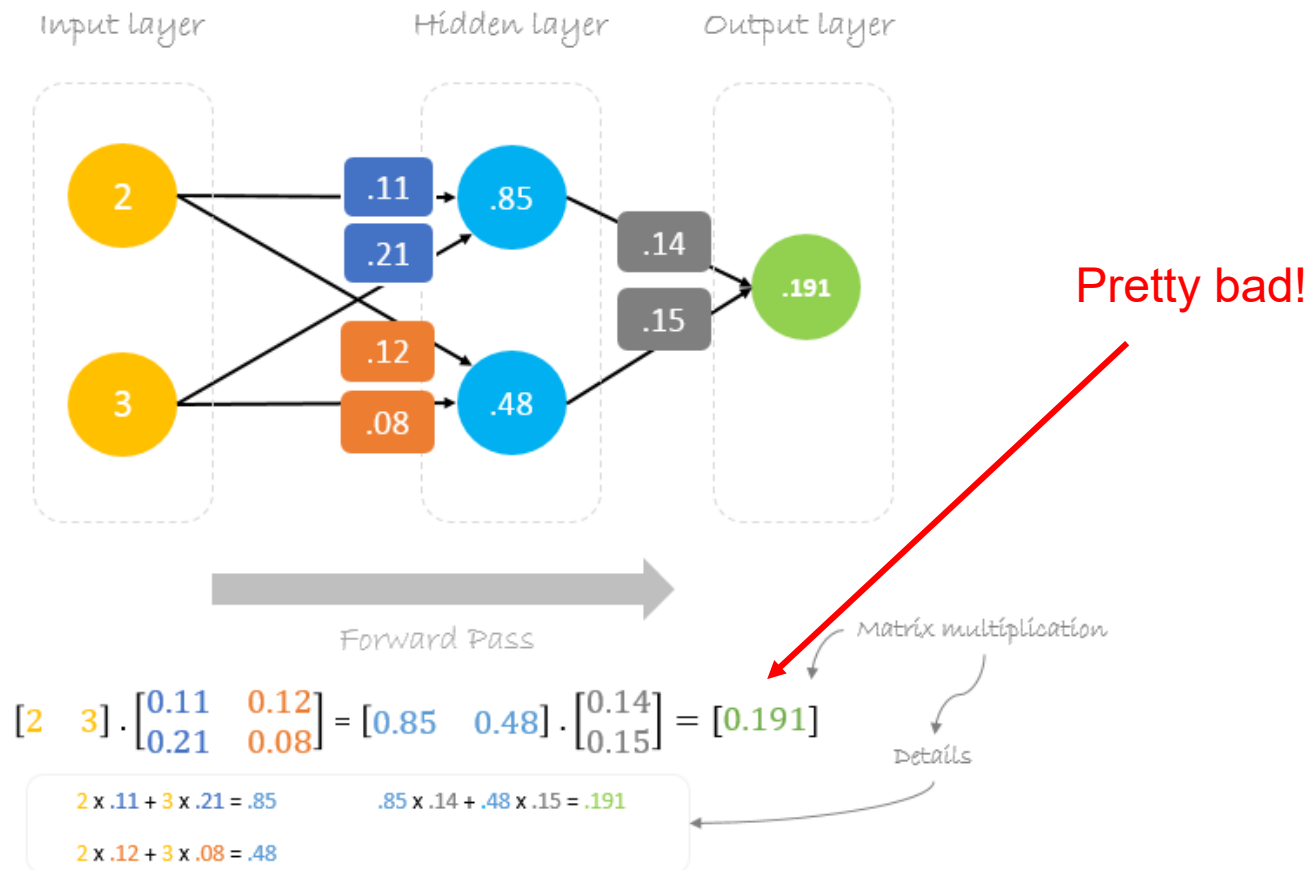
Backpropagation

Forward pass



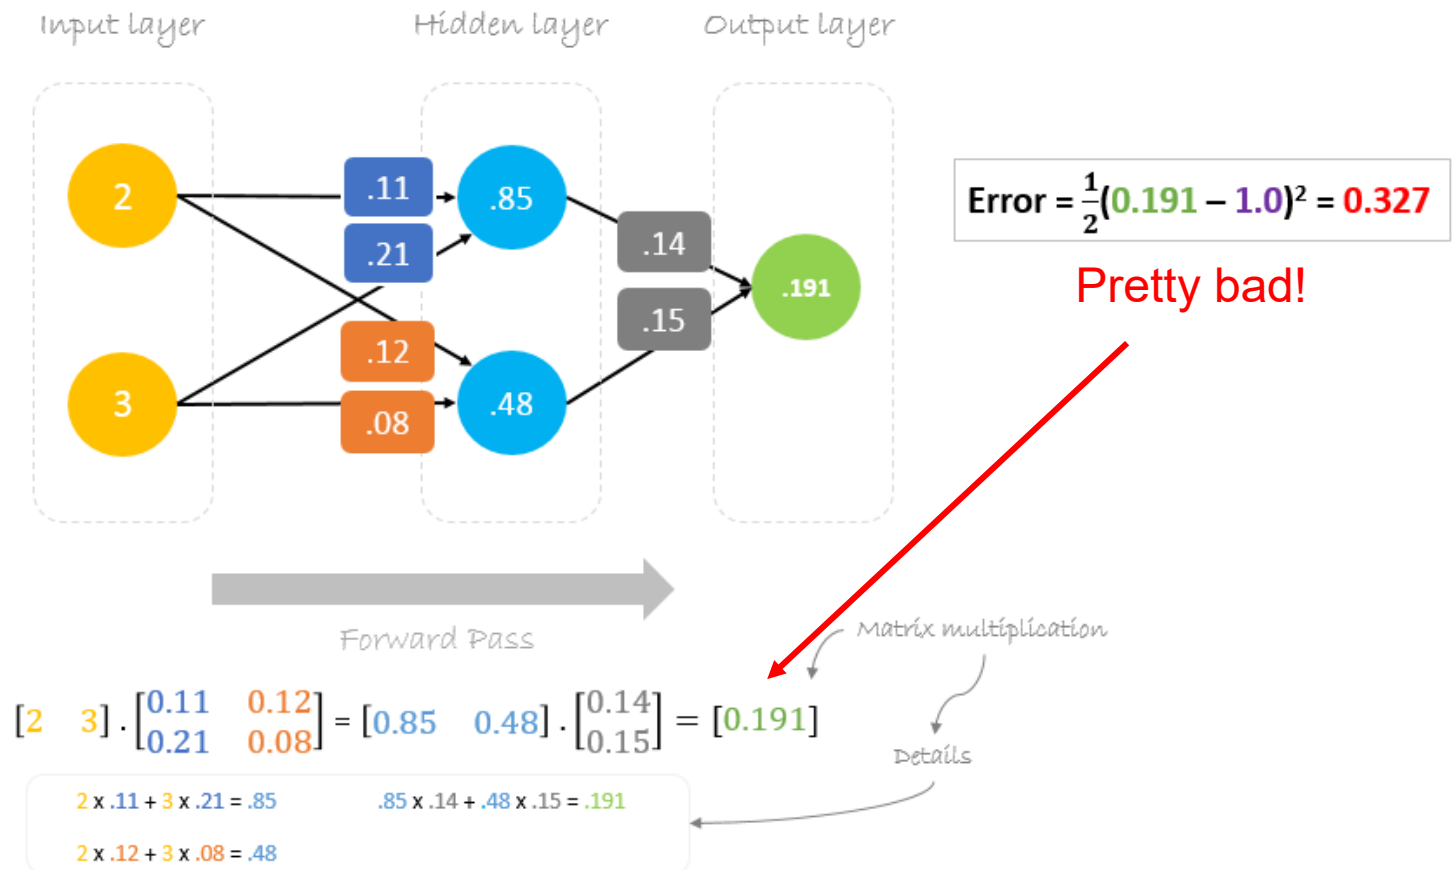
Backpropagation

Forward pass



Backpropagation

Forward pass



Backpropagation

Backpropagate for update

$$*W_x = W_x - a \left(\frac{\partial \text{Error}}{\partial W_x} \right)$$

Diagram illustrating the weight update formula during backpropagation:

- $*W_x$: New weight
- W_x : Old weight
- a : Learning rate
- $\left(\frac{\partial \text{Error}}{\partial W_x} \right)$: Derivative of Error with respect to weight

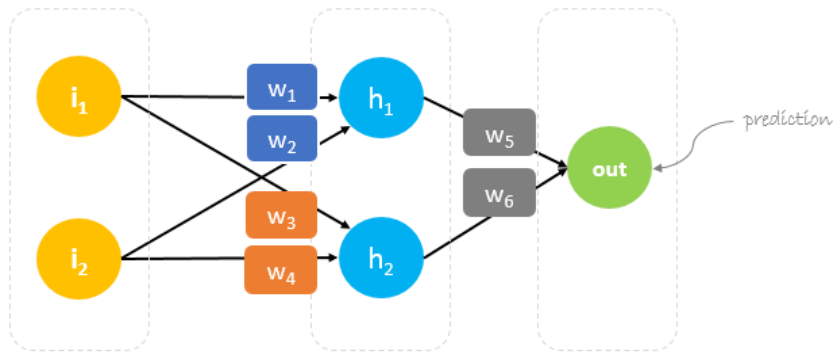
Backpropagation

Backpropagate for update

$$*W_x = W_x - \alpha \left(\frac{\partial \text{Error}}{\partial W_x} \right)$$

Annotations:

- Old weight: points to W_x
- Derivative of Error with respect to weight: points to $\left(\frac{\partial \text{Error}}{\partial W_x} \right)$
- New weight: points to $*W_x$
- Learning rate: points to α



$$\begin{aligned} & \text{out} \\ & \downarrow \\ & (h_1) w_5 + (h_2) w_6 \\ & \downarrow \\ & (i_1 w_1 + i_2 w_2) w_5 + (i_1 w_3 + i_2 w_4) w_6 \end{aligned}$$

Intermediate calculations:

- $h_1 = i_1 w_1 + i_2 w_2$
- $h_2 = i_1 w_3 + i_2 w_4$

Backpropagation

out
↓

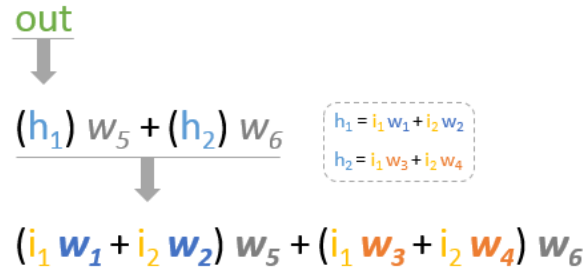
$$\frac{(h_1) w_5 + (h_2) w_6}{\downarrow}$$

$(i_1 w_1 + i_2 w_2) w_5 + (i_1 w_3 + i_2 w_4) w_6$

$$\begin{aligned} h_1 &= i_1 w_1 + i_2 w_2 \\ h_2 &= i_1 w_3 + i_2 w_4 \end{aligned}$$

$$\frac{\partial Error}{\partial y} = (y - y^*)$$

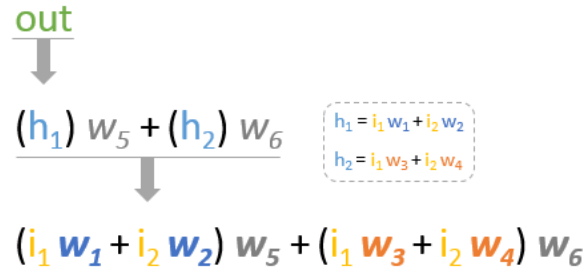
Backpropagation



$$\frac{\partial Error}{\partial y} = (y - y^*)$$

$$\frac{\partial Error}{\partial W_6} = \frac{\partial Error}{\partial y} \frac{\partial y}{\partial W_6} = (y - y^*) h_2$$

Backpropagation

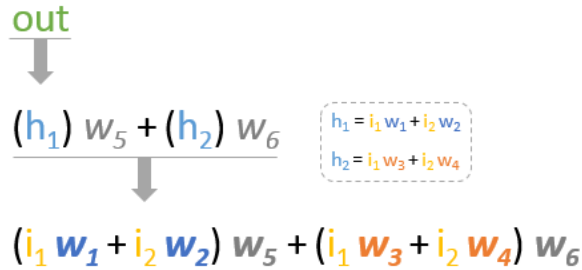


$$\frac{\partial Error}{\partial y} = (y - y^*)$$

$$\frac{\partial Error}{\partial W_6} = \frac{\partial Error}{\partial y} \frac{\partial y}{\partial W_6} = (y - y^*) h_2$$

$$\frac{\partial Error}{\partial W_5} = \frac{\partial Error}{\partial y} \frac{\partial y}{\partial W_5} = (y - y^*) h_1$$

Backpropagation



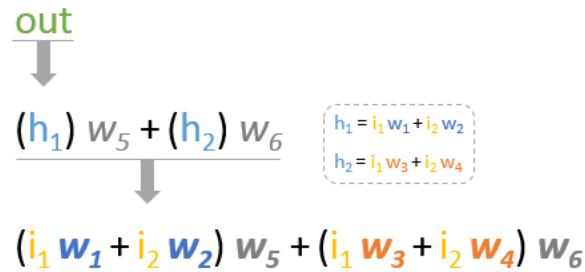
$$\frac{\partial Error}{\partial y} = (y - y^*)$$

$$\frac{\partial Error}{\partial W_6} = \frac{\partial Error}{\partial y} \frac{\partial y}{\partial W_6} = (y - y^*) h_2$$

$$\frac{\partial Error}{\partial W_5} = \frac{\partial Error}{\partial y} \frac{\partial y}{\partial W_5} = (y - y^*) h_1$$

$$\frac{\partial Error}{\partial W_1} = \frac{\partial Error}{\partial y} \frac{\partial y}{\partial h_1} \frac{\partial h_1}{\partial W_1} = (y - y^*) W_5 i_1$$

Backpropagation



$$\frac{\partial Error}{\partial y} = (y - y^*)$$

$$\frac{\partial Error}{\partial W_6} = \frac{\partial Error}{\partial y} \frac{\partial y}{\partial W_6} = (y - y^*) h_2$$

$$\frac{\partial Error}{\partial W_5} = \frac{\partial Error}{\partial y} \frac{\partial y}{\partial W_5} = (y - y^*) h_1$$

$$\frac{\partial Error}{\partial W_1} = \frac{\partial Error}{\partial y} \frac{\partial y}{\partial h_1} \frac{\partial h_1}{\partial W_1} = (y - y^*) W_5 i_1$$

etc

Backpropagation

out
↓

$$\frac{(h_1) w_5 + (h_2) w_6}{\downarrow}$$

$$(i_1 w_1 + i_2 w_2) w_5 + (i_1 w_3 + i_2 w_4) w_6$$

$$\begin{aligned} h_1 &= i_1 w_1 + i_2 w_2 \\ h_2 &= i_1 w_3 + i_2 w_4 \end{aligned}$$

For learning rate a , $\Delta = (y - y^*)$

we have the following updates

Backpropagation

out
↓

$$\frac{(h_1) w_5 + (h_2) w_6}{(i_1 w_1 + i_2 w_2) w_5 + (i_1 w_3 + i_2 w_4) w_6}$$

$$\begin{aligned} h_1 &= i_1 w_1 + i_2 w_2 \\ h_2 &= i_1 w_3 + i_2 w_4 \end{aligned}$$

For learning rate a , $\Delta = (y - y^*)$

we have the following updates

$$\begin{bmatrix} w_5 \\ w_6 \end{bmatrix} = \begin{bmatrix} w_5 \\ w_6 \end{bmatrix} - a \Delta \begin{bmatrix} h_1 \\ h_2 \end{bmatrix} = \begin{bmatrix} w_5 \\ w_6 \end{bmatrix} - \begin{bmatrix} a h_1 \Delta \\ a h_2 \Delta \end{bmatrix}$$

$$\begin{bmatrix} w_1 & w_3 \\ w_2 & w_4 \end{bmatrix} = \begin{bmatrix} w_1 & w_3 \\ w_2 & w_4 \end{bmatrix} - a \Delta \begin{bmatrix} i_1 \\ i_2 \end{bmatrix} \cdot [w_5 \quad w_6] = \begin{bmatrix} w_1 & w_3 \\ w_2 & w_4 \end{bmatrix} - \begin{bmatrix} a i_1 \Delta w_5 & a i_1 \Delta w_6 \\ a i_2 \Delta w_5 & a i_2 \Delta w_6 \end{bmatrix}$$

Backpropagation

Let $a = 0.05$

we know $\Delta = (y - y^*) = 0.191 - 1 = -0.809$

Thus

$$\begin{bmatrix} w_5 \\ w_6 \end{bmatrix} = \begin{bmatrix} 0.14 \\ 0.15 \end{bmatrix} - 0.05(-0.809) \begin{bmatrix} 0.85 \\ 0.48 \end{bmatrix} = \begin{bmatrix} 0.14 \\ 0.15 \end{bmatrix} - \begin{bmatrix} -0.034 \\ -0.019 \end{bmatrix} = \begin{bmatrix} 0.17 \\ 0.17 \end{bmatrix}$$

$$\begin{bmatrix} w_1 & w_3 \\ w_2 & w_4 \end{bmatrix} = \begin{bmatrix} .11 & .12 \\ .21 & .08 \end{bmatrix} - 0.05(-0.809) \begin{bmatrix} 2 \\ 3 \end{bmatrix} \cdot [0.14 \quad 0.15] = \begin{bmatrix} .11 & .12 \\ .21 & .08 \end{bmatrix} - \begin{bmatrix} -0.011 & -0.012 \\ -0.017 & -0.018 \end{bmatrix} = \begin{bmatrix} .12 & .13 \\ .23 & .10 \end{bmatrix}$$

Backpropagation

Let $a = 0.05$

we know $\Delta = (y - y^*) = 0.191 - 1 = -0.809$

Thus

$$\begin{bmatrix} w_5 \\ w_6 \end{bmatrix} = \begin{bmatrix} 0.14 \\ 0.15 \end{bmatrix} - 0.05(-0.809) \begin{bmatrix} 0.85 \\ 0.48 \end{bmatrix} = \begin{bmatrix} 0.14 \\ 0.15 \end{bmatrix} - \begin{bmatrix} -0.034 \\ -0.019 \end{bmatrix} = \begin{bmatrix} 0.17 \\ 0.17 \end{bmatrix}$$

$$\begin{bmatrix} w_1 & w_3 \\ w_2 & w_4 \end{bmatrix} = \begin{bmatrix} .11 & .12 \\ .21 & .08 \end{bmatrix} - 0.05(-0.809) \begin{bmatrix} 2 \\ 3 \end{bmatrix} \cdot \begin{bmatrix} 0.14 & 0.15 \end{bmatrix} = \begin{bmatrix} .11 & .12 \\ .21 & .08 \end{bmatrix} - \begin{bmatrix} -0.011 & -0.012 \\ -0.017 & -0.018 \end{bmatrix} = \begin{bmatrix} .12 & .13 \\ .23 & .10 \end{bmatrix}$$

Sanity check: With the new weights do a forward pass

The new prediction becomes $y = 0.26 \rightarrow$ better!

Backpropagation

Let $a = 0.05$

we know $\Delta = (y - y^*) = 0.191 - 1 = -0.809$

Thus

$$\begin{bmatrix} w_5 \\ w_6 \end{bmatrix} = \begin{bmatrix} 0.14 \\ 0.15 \end{bmatrix} - 0.05(-0.809) \begin{bmatrix} 0.85 \\ 0.48 \end{bmatrix} = \begin{bmatrix} 0.14 \\ 0.15 \end{bmatrix} - \begin{bmatrix} -0.034 \\ -0.019 \end{bmatrix} = \begin{bmatrix} 0.17 \\ 0.17 \end{bmatrix}$$

$$\begin{bmatrix} w_1 & w_3 \\ w_2 & w_4 \end{bmatrix} = \begin{bmatrix} .11 & .12 \\ .21 & .08 \end{bmatrix} - 0.05(-0.809) \begin{bmatrix} 2 \\ 3 \end{bmatrix} \cdot \begin{bmatrix} 0.14 & 0.15 \end{bmatrix} = \begin{bmatrix} .11 & .12 \\ .21 & .08 \end{bmatrix} - \begin{bmatrix} -0.011 & -0.012 \\ -0.017 & -0.018 \end{bmatrix} = \begin{bmatrix} .12 & .13 \\ .23 & .10 \end{bmatrix}$$

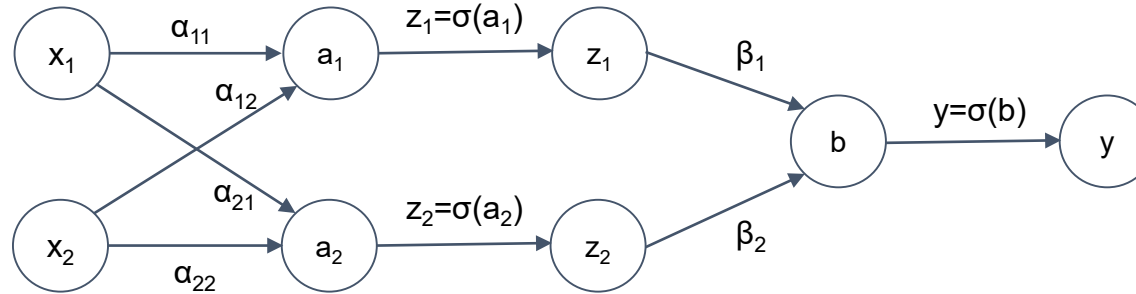
Sanity check: With the new weights do a forward pass

REPEAT!

The new prediction becomes $y = 0.26 \rightarrow$ better!

Backpropagation

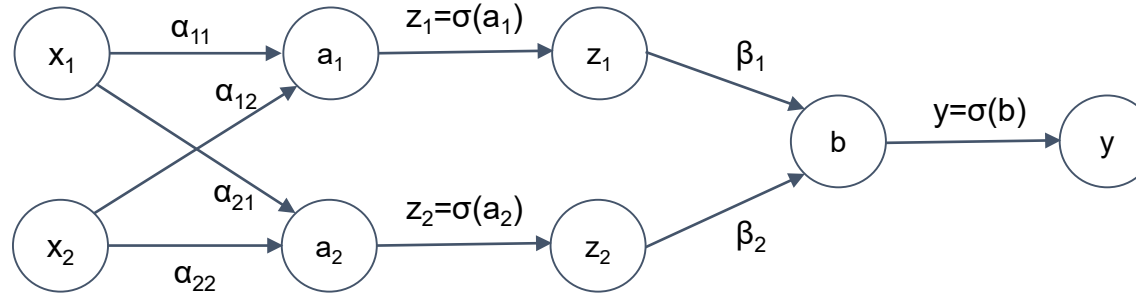
- More realistic example



- Let the activation σ be sigmoid $\sigma(z) = \frac{1}{1 + e^{-z}}$
- Need to learn $\alpha_{11}, \alpha_{12}, \alpha_{21}, \alpha_{22}, \beta_1, \beta_2$

Backpropagation

- More realistic example

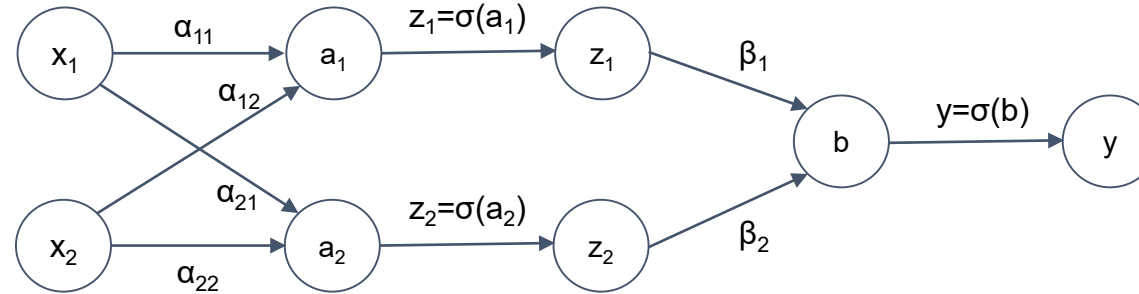


$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- We will consider logistic regression
 - Models probability for a binary classification problem
 - Outputs are 0 and 1

Backpropagation

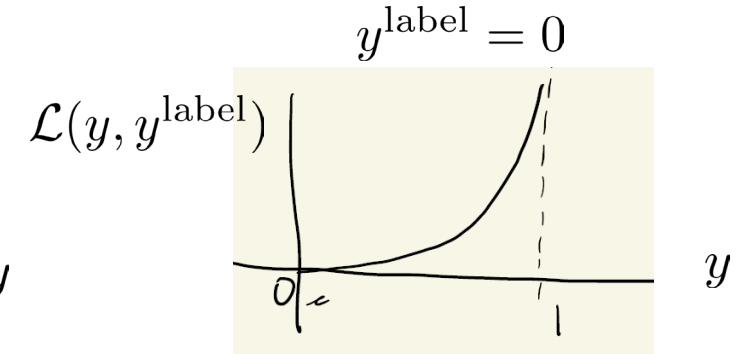
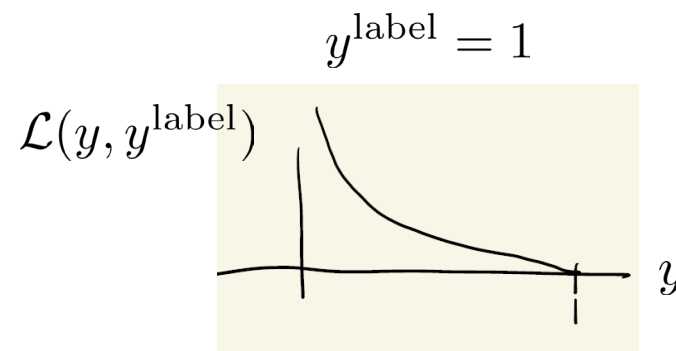
- More realistic example



$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

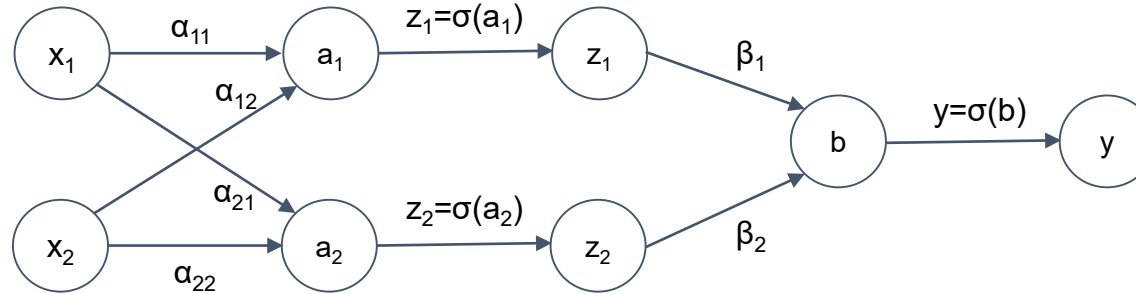
- We will consider logistic regression

$$\mathcal{L}(y, y^{\text{label}}) = \begin{cases} -\log(y) & \text{if } y^{\text{label}} = 1 \\ -\log(1 - y) & \text{if } y^{\text{label}} = 0 \end{cases}$$



Backpropagation

- More realistic example



$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- We will consider logistic regression

$$\mathcal{L}(y, y^{\text{label}}) = \begin{cases} -\log(y) & \text{if } y^{\text{label}} = 1 \\ -\log(1 - y) & \text{if } y^{\text{label}} = 0 \end{cases}$$

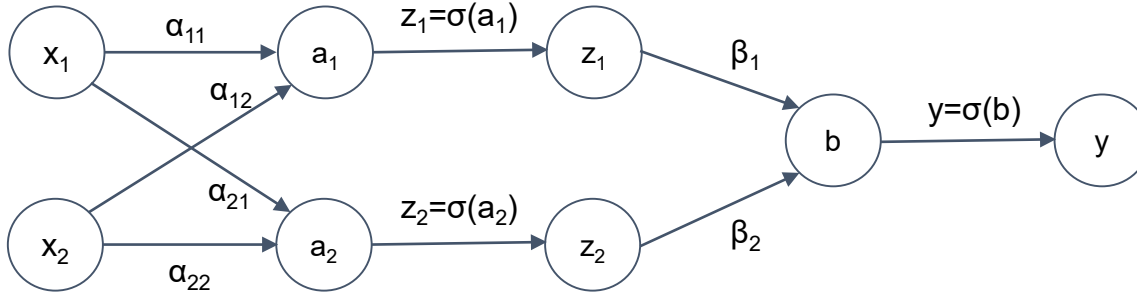
or

$$\mathcal{L}(y, y^{\text{label}}) = -y^{\text{label}} \log(y) - (1 - y^{\text{label}}) \log(1 - y)$$

- Again, we need to backpropagate

Backpropagation

- More realistic example



Forward pass

$$y = \frac{1}{1 + e^{-b}}$$

$$b = \beta_1 z_1 + \beta_2 z_2$$

w.r.t. layer
parameters

w.r.t. input to
backpropagate

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$$\mathcal{L}(y, y^{\text{label}}) = -y^{\text{label}} \log(y) - (1 - y^{\text{label}}) \log(1 - y)$$

$$\frac{\partial \mathcal{L}}{\partial y} = -\frac{y^{\text{label}}}{y} + \frac{1 - y^{\text{label}}}{1 - y}$$

Backward pass

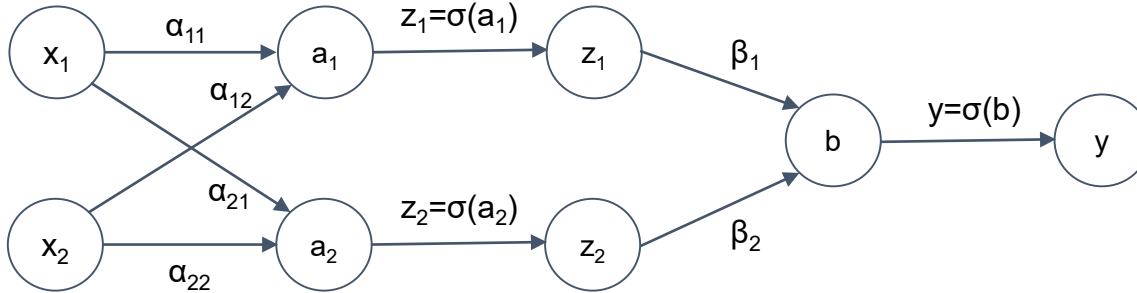
$$\frac{\partial \mathcal{L}}{\partial b} = \frac{\partial \mathcal{L}}{\partial y} \frac{\partial y}{\partial b} \quad \frac{\partial y}{\partial b} = \frac{e^{-b}}{(1 + e^{-b})^2} = y(1 - y)$$

$$\frac{\partial \mathcal{L}}{\partial \beta_i} = \frac{\partial \mathcal{L}}{\partial b} \frac{\partial b}{\partial \beta_i} \quad \frac{\partial b}{\partial \beta_i} = z_i$$

$$\frac{\partial \mathcal{L}}{\partial z_i} = \frac{\partial \mathcal{L}}{\partial b} \frac{\partial b}{\partial z_i} \quad \frac{\partial b}{\partial z_i} = \beta_i$$

Backpropagation

- More realistic example



Forward pass

$$z_j = \frac{1}{1 + e^{-a_j}}$$

$$a_j = \alpha_{j1}x_1 + \alpha_{j2}x_2$$

w.r.t. layer
parameters

w.r.t. input to
backpropagate

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$$\mathcal{L}(y, y^{\text{label}}) = -y^{\text{label}} \log(y) - (1 - y^{\text{label}}) \log(1 - y)$$

$$\frac{\partial \mathcal{L}}{\partial y} = -\frac{y^{\text{label}}}{y} + \frac{1 - y^{\text{label}}}{1 - y}$$

Backward pass

$$\frac{\partial \mathcal{L}}{\partial a_j} = \frac{\partial \mathcal{L}}{\partial z_j} \frac{\partial z_j}{\partial a_j} \quad \frac{\partial z_j}{\partial a_j} = z_j(1 - z_j)$$

$$\frac{\partial \mathcal{L}}{\partial \alpha_{jk}} = \frac{\partial \mathcal{L}}{\partial a_j} \frac{\partial a_j}{\partial \alpha_{jk}} \quad \frac{\partial a_j}{\partial \alpha_{jk}} = x_k$$

$$\frac{\partial \mathcal{L}}{\partial x_k} = \frac{\partial \mathcal{L}}{\partial a_1} \frac{\partial a_1}{\partial x_k} + \frac{\partial \mathcal{L}}{\partial a_2} \frac{\partial a_2}{\partial x_k}$$

etc

In Practice

- Most operators will be vector valued
 - Derivatives replaced with Jacobian matrices
- The software will do the differentiation for you
- Other derivatives

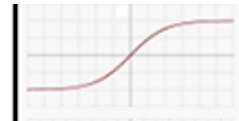
Rectified
Linear Unit
(ReLU)



$$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$$

$$f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$$






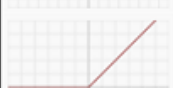



Tanh



$$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$$

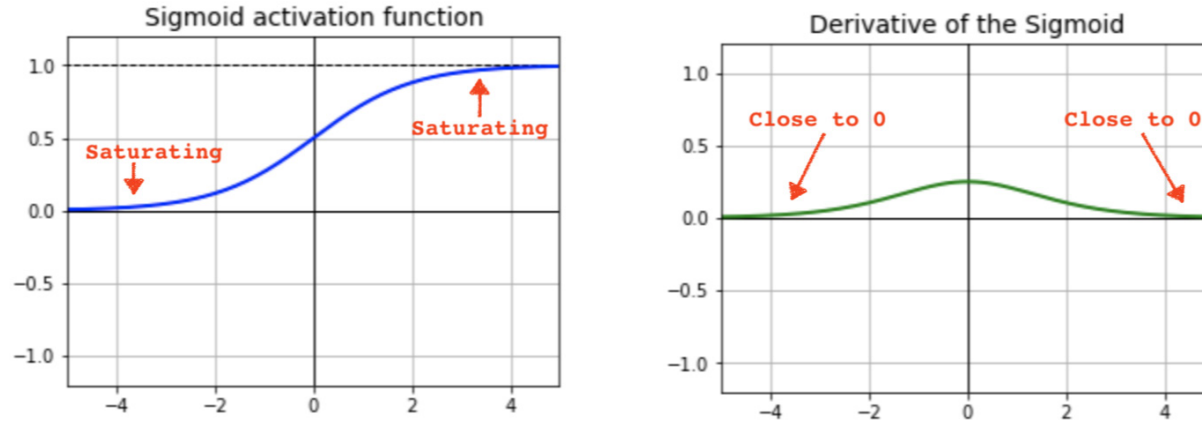
$$f'(x) = 1 - f(x)^2$$

Derivatives of Activation Functions

Name	Plot	Equation	Derivative
Identity		$f(x) = x$	$f'(x) = 1$
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x \neq 0 \\ ? & \text{for } x = 0 \end{cases}$
Logistic (a.k.a. Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$	$f'(x) = f(x)(1 - f(x))$
Tanh		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$	$f'(x) = 1 - f(x)^2$
ArcTan		$f(x) = \tan^{-1}(x)$	$f'(x) = \frac{1}{x^2 + 1}$
Rectified Linear Unit (ReLU)		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Parameteric Rectified Linear Unit (PReLU)		$f(x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Exponential Linear Unit (ELU)		$f(x) = \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} f(x) + \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
SoftPlus		$f(x) = \log_e(1 + e^x)$	$f'(x) = \frac{1}{1 + e^{-x}}$

In Practice

- What do these mean for multi-layer networks & backpropagation?



- If backpropagating over many layers, these will get multiplied
 - Small derivatives (exponentially decreasing with distance from last layer) → Numerical issues
 - Called “vanishing gradients”
 - Same problem for tanh
- For ReLU → Derivative is either 0 or 1
 - Solves this issue → Thus it's very popular

In Practice

$$\arg \min_{\theta} \sum_{k=1}^n \mathcal{L}(f_{\theta}(x_k), y_k)$$

– We can compute the gradient over

- Entire training database

$$\theta = \theta - \eta \sum_{k=1}^n \nabla_{\theta} \mathcal{L}(f_{\theta}(x_k), y_k)$$

- Single data point

$$\theta = \theta - \eta \nabla_{\theta} \mathcal{L}(f_{\theta}(x_k), y_k)$$

- A smaller subset (mini-batch)

$$\theta = \theta - \eta \sum_{k=k_0+1}^{k_0+n_0} \nabla_{\theta} \mathcal{L}(f_{\theta}(x_k), y_k)$$

← This is what is usually done

Stochastic Gradient Descent (SGD)

– Stochastic gradient descent (SGD)

- Randomly picks a mini-batch and uses this for gradient calculations and updates

$$\theta = \theta - \eta \sum_{k \in \mathcal{K}_j} \nabla_{\theta} \mathcal{L}(f_{\theta}(x_k), y_k) \quad \mathcal{K}_j \text{ randomly chosen index set of size } n_0$$

- Note the random index set itself is indexed over j
 - j has to vary between 1 to $\text{round}(n/n_0)$ to cover the entire training dataset
 - Note mini-batches are chosen in a way to cover the entire dataset (sampling without replacement)
- Two different terms:
 - Iteration: Each gradient update per mini-batch
$$\theta = \theta - \eta \sum_{k \in \mathcal{K}_j} \nabla_{\theta} \mathcal{L}(f_{\theta}(x_k), y_k)$$
 - Epoch: When the entire dataset has been passed forward & backward once
 - n/n_0 iterations needed in an epoch

Stochastic Gradient Descent (SGD)

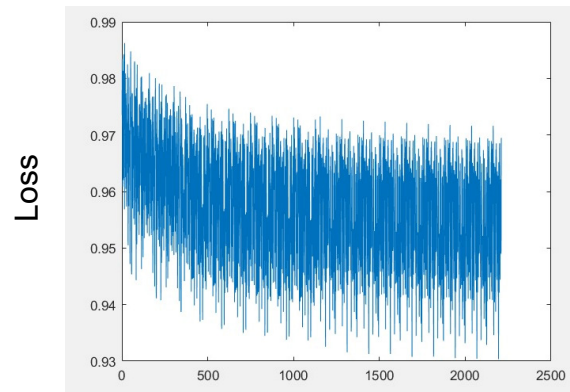
– Stochastic gradient descent (SGD)

- Randomly picks a mini-batch and uses this for gradient calculations and updates

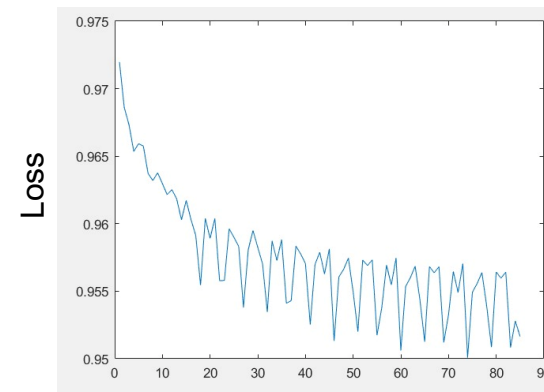
$$\theta = \theta - \eta \sum_{k \in \mathcal{K}_j} \nabla_{\theta} \mathcal{L}(f_{\theta}(x_k), y_k)$$

\mathcal{K}_j randomly chosen index set of size n_0

- Note the random index set itself is indexed over j
 - j has to vary between 1 to $\text{round}(n/n_0)$ to cover the entire training dataset
 - Note mini-batches are chosen in a way to cover the entire dataset (sampling without replacement)
- Iteration vs. epoch



Iteration #



Epoch #

SGD

- Stochastic gradient descent (SGD)
 - In practice with images batch sizes of 32-64 (or higher now) are common
 - Depends on application!
 - η is called *learning rate* in deep learning applications (we used to call it step size)
 - SGD is much faster & more memory-efficient than GD
 - Can use optimized matrix operations in deep learning libraries
 - Reduces the variance of the parameter updates due to random selection
 - Can escape suboptimal local minima easier than GD
 - Converges faster than GD

More on SGD

– Stochastic gradient descent (SGD)

- Many many variants exist...
- One important concept is momentum
 - To avoid getting trapped in bad local minima
 - If the optimization surface is steeper than the other SGD will oscillate
 - Momentum dampens oscillations



Image 2: SGD without momentum



Image 3: SGD with momentum

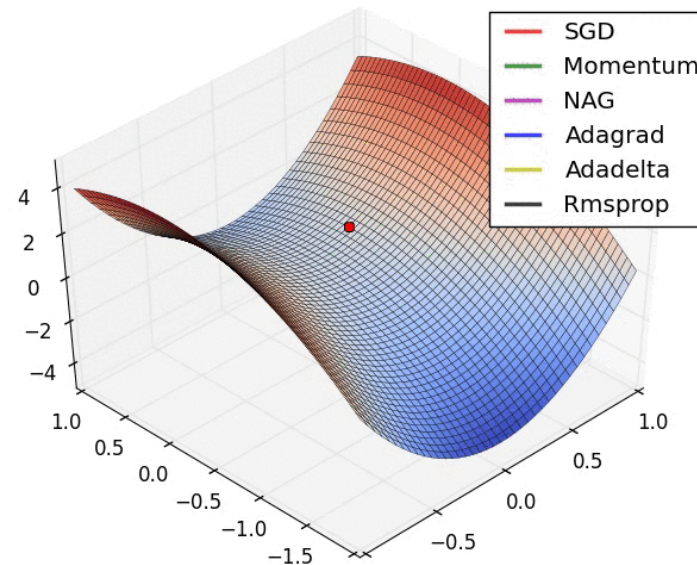
- Does this by adding a fraction of the update vector of the past time step to the current update vector

$$\mathbf{v}^{(t)} = \gamma \mathbf{v}^{(t-1)} + \eta \sum_{k \in \mathcal{K}_j} \nabla_{\boldsymbol{\theta}} \mathcal{L}(f_{\boldsymbol{\theta}}(x_k), y_k)$$

$$\boldsymbol{\theta}^{(t)} = \boldsymbol{\theta}^{(t-1)} - \mathbf{v}^{(t)}$$

More on SGD

- Stochastic gradient descent (SGD)
 - Many many variants exist...
 - Other important concept: learning rate “schedules”
 - Hard to choose a good learning rate (very important hyperparameter in practice!)
 - High learning rate is good in early epochs, and low learning rate is better at later epochs
 - Also may need different learning rates for different parameters
 - AdaGrad, RMSProp, Adam ...



This is a gif

<http://ruder.io/optimizing-gradient-descent/>

In Practice

- How do we initialize the weights in θ ?
 - Mostly randomly with zero-mean
 - Either constant variance across layers
 - Or normalized with respect to number of channels/layers (Xavier initialization)
 - Most networks now have fixed number of channels/layer \rightarrow approximately leads to the first setting

- Should we regularize the weights?

$$\arg \min_{\theta} \sum_{k=1}^n \mathcal{L}(f_{\theta}(x_k), y_k) + \mathcal{R}(\theta)$$

- We can replicate what we have done before (e.g. l_p norm regularizers)
- We will see there are other methods for regularization
 - e.g. implicitly in the SGD via the random selection of indices

Recap

- Backpropagation
 - Concept & examples for clarification
 - In practice: All done automatically in the learning framework
- Practical training points
 - Optimization algorithms
 - Hyperparameters/initialization
 - Regularization
- Course announcement
 - Office hours shortened (11:00-11:45) via [zoom](#) tomorrow