

EE 5561: Image Processing and Applications

Lecture 19

Mehmet Akçakaya

Recap of Last Lecture

- Convolutional neural networks
 - Why convolutions instead of fully-connected layers?
 - Sparse weights
 - Parameter sharing
 - Translation equivariance
 - Works with inputs of different sizes
 - Variations on convolutions & visualizing convolutional layers
 - Other CNN components
 - Pooling layers
 - Normalization layers
 - Activation layers (from earlier)

Regularization in Neural Networks

- We talked about loss functions in training

$$\mathcal{L}(f_{\theta}(X), Y) \triangleq J(\theta; X, Y)$$

loss function network parameters input output

- We minimized $J(\theta; X, Y)$ with respect to θ , when training the neural network
- We might do additional changes to $J(\theta; X, Y)$ to avoid overfitting or limit the capacity of the model
 - This is the concept of regularization (which we have spent quite a bit of time on)
- Idea: Instead minimize

$$J(\theta; X, Y) + \alpha \mathcal{R}(\theta)$$

still defined over the database positive hyperparameter regularization term

Norm Regularizers

- What we know best!
- Generally: l_p norm regularizers, such as

- l_1 norm: $||\boldsymbol{\theta}||_1 = \sum_k |\theta_k|$

- l_2 norm: $||\boldsymbol{\theta}||_2 = \left(\sum_k |\theta_k|^2 \right)^{\frac{1}{2}}$

- l_∞ norm: $||\boldsymbol{\theta}||_\infty = \max_k |\theta_k|$

- l_p norm: $||\boldsymbol{\theta}||_p = \left(\sum_k |\theta_k|^p \right)^{\frac{1}{p}} \quad p \geq 1$

Norm Regularizers

- We know from before that l_2 norm is the computationally easiest to work with

- Ridge regression (if MSE loss and linear model), a.k.a. Tikhonov regularization
- For training, we now minimize

$$J(\boldsymbol{\theta}; X, Y) + \frac{\alpha}{2} \|\boldsymbol{\theta}\|_2^2 = J(\boldsymbol{\theta}; X, Y) + \frac{\alpha}{2} \boldsymbol{\theta}^T \boldsymbol{\theta}$$

- Note its gradient is given as

$$\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}; X, Y) + \alpha \boldsymbol{\theta}$$

- Thus the gradient update would be

$$\begin{aligned} \boldsymbol{\theta} &\leftarrow \boldsymbol{\theta} - \epsilon (\alpha \boldsymbol{\theta} + \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}; X, Y)) \\ &= (1 - \epsilon \alpha) \boldsymbol{\theta} - \epsilon \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}; X, Y) \end{aligned}$$

↑
weight decay

Norm Regularizers

– We also know that l_1 norm regularizer is important

- Related to sparsity \rightarrow feature selection
- Robust to outliers
- For training, we now minimize

$$J(\boldsymbol{\theta}; X, Y) + \alpha \|\boldsymbol{\theta}\|_1$$


- Note its gradient is given as

$$\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}; X, Y) + \alpha \cdot \text{sign}(\boldsymbol{\theta})$$


- Computationally more difficult to work with

Norm Regularizers

- What we know from basic statistical image processing applies here too
 - l_2 norm regularization ~ Gaussian prior on weights



- l_1 norm regularization ~ sparsity-promoting prior on weights



Regularization in Neural Networks

- Norm regularizers are theoretically easy to understand
 - Well-defined behavior on how it changes the optimization
 - But usually computationally hard to work with
 - Except l_2 norm, which leads to a very simple regularizer (weight decay)
- Thus, there are multiple non-norm regularizers used in neural networks


Data Augmentation

- Recall that our end goal is to have the best generalizability possible for our trained network
- Best way to improve generalizability: Get more data
- If we can't get more “real” data, we can “fake” the data
- Called data augmentation




Data Augmentation


Crop




Scale




Flip




Rotate



Translate



Add noise



Why these?

- In real life, databases of images are taken in a limited set of conditions
- In applications, we want to do well in a variety of conditions (different orientation, scaling, brightness, etc)

Data Augmentation

– Augmentation may help even if we have lots of data?

- May increase the amount of relevant data
- Consider the hypothetical scenario:
 - Database of images with brand A (Ford) and brand B (Chevrolet)
 - All brand A cars face left, all brand B cars face right
 - Train network with a huge database
 - At test time




- Network will pick brand B (right facing), even though it is brand A
- Because the neural network learns features that distinguish two classes
- The most relevant feature here is right-left facing
- Augmentation would help here!

Data Augmentation

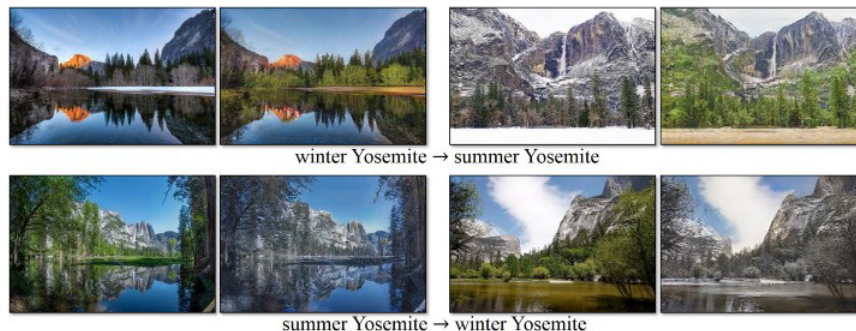
– Practical points:

- Some spatial transformation (e.g. rotation/translation) may lead to areas out of the field-of-view (or black regions)
- Need to decide what to do with these
 - Idea ~boundary processing (e.g. reflect, symmetric, wrap, or do nothing)




– More advanced augmentation methods exist

- Add noise to intermediate layers or output (“label smoothing”)
- Generative models (e.g. convert summer background to winter background) [Later]



Early Stopping

- Stop the training when validation set loss goes up



Early Stopping

– Why regularization?

- We are effectively restricting the optimization to a “small” volume of the parameter space near the initial parameter value
- Taylor series expansion around the optimal solution θ^*

$$J(\theta; X, Y) = J(\theta^*; X, Y) + \frac{1}{2}(\theta - \theta^*)^T \mathbf{H}(\theta - \theta^*)$$

Note $\nabla_{\theta} J(\theta; X, Y)|_{\theta=\theta^*} = 0$ by optimality of θ^*


Thus $\nabla_{\theta} J(\theta; X, Y) = \mathbf{H}(\theta - \theta^*)$

- At iteration t we have the gradient update

$$\theta^{(t)} = \theta^{(t-1)} - \epsilon \mathbf{H}(\theta^{(t-1)} - \theta^*)$$

$$\theta^{(t)} - \theta^* = (\mathbf{I} - \epsilon \mathbf{H})(\theta^{(t-1)} - \theta^*) = (\mathbf{I} - \epsilon \mathbf{H})^t(\theta^{(0)} - \theta^*)$$

- If loss is MSE, one can show this is similar to l_2 regularization with $t \approx \frac{1}{\epsilon \alpha}$




Dropout Regularization


- One of the most important regularization methods for neural networks
- Co-adaptation: When all weights are learned together, some of the connections have more predictive capability than others
 - Strong weights become stronger, weak weights become weaker
 - So only a fraction of nodes are actually trained
 - Problem: Expanding the network size does not help → accuracy remains limited
- Idea: Randomly drop nodes from computation graph throughout training

Dropout Regularization

- Idea: Randomly drop nodes from computation graph throughout training
 - Output turned to zero and not participating in back-propagation
 - ~ Training multiple networks with different compositions & “averaging” their results



(a) Standard Neural Net




(b) After applying dropout.

- Note this is not “ensemble learning”, which trains multiple models and averages those → too expensive in practice
- Dropout drops random neurons with *every* pass through the network
 - Essentially sees a “new network” at every pass
 - Weights are still shared between these “networks” (unlike ensemble learning)

Dropout Regularization

- Idea: Randomly drop nodes from computation graph throughout training
 - Set a dropout rate, $p < 1$, for each layer (probability that a unit is dropped)
 - This is a random variable, Bernoulli(p)




Dropout Regularization

– Probabilistic Formulation

Normal Feedforward Net


$$\begin{aligned}z_i^{(l+1)} &= \mathbf{w}_i^{(l+1)} \mathbf{y}^l + b_i^{(l+1)} \\ y_i^{(l+1)} &= f(z_i^{(l+1)}),\end{aligned}$$



(a) Standard network

With Dropout



$$\begin{aligned}r_j^{(l)} &\sim \text{Bernoulli}(p), \\ \tilde{\mathbf{y}}^{(l)} &= \mathbf{r}^{(l)} * \mathbf{y}^{(l)}, \\ z_i^{(l+1)} &= \mathbf{w}_i^{(l+1)} \tilde{\mathbf{y}}^l + b_i^{(l+1)} \\ y_i^{(l+1)} &= f(z_i^{(l+1)}).\end{aligned}$$



(b) Dropout network


Dropout Regularization

- At test time:
 - Want to use the full network
 - But now the neurons will receive more connections (& more activations), e.g.
 - If $p = 0.5$, then on average we will drop 2 out of 4 connections in a layer during training
 - But at testing, they will all be active \rightarrow The response they produce will be large by ~ 2 -fold
 - There is a correction added to compensate for this effect
 - Scale the weights by p at test time



Dropout Regularization

- On MNIST



- Variants exist

- e.g. replace Bernoulli RV with Gaussian RV $N(1, \sigma^2)$
- Advantage: Expected value is 1 \rightarrow So no change during inference time

Types of Learning Problems

- A lot of techniques named “XX learning”
 - Here we will try to cover some of these relevant for our tasks
 - We will split them into two: Learning tasks & learning methods
 - Learning tasks: Supervised learning, unsupervised learning, semi-supervised learning, self-supervised learning, ...
 - Learning methods: Active learning, transfer learning, federated learning, ...
 - We have been looking at supervised learning so far
 - Matched input & label data, train using these matched information
 - We will first look at “learning methods” in the supervised setting
 - Then we will revisit learning tasks

Learning Techniques

- These are the methodologies with which you learn a neural network for an application
- Active learning
 - Strategy to identify which specific examples/labels in our training data can best improve model performance
 - Simple example: Oracle (human-in-the-loop) identifies important examples
 - More advanced techniques automate this (e.g. via uncertainty quantification)
- Ensemble learning
 - Train ≥ 2 models on some data, then combine predictions
 - Idea: Ensemble performance $>$ individual mode

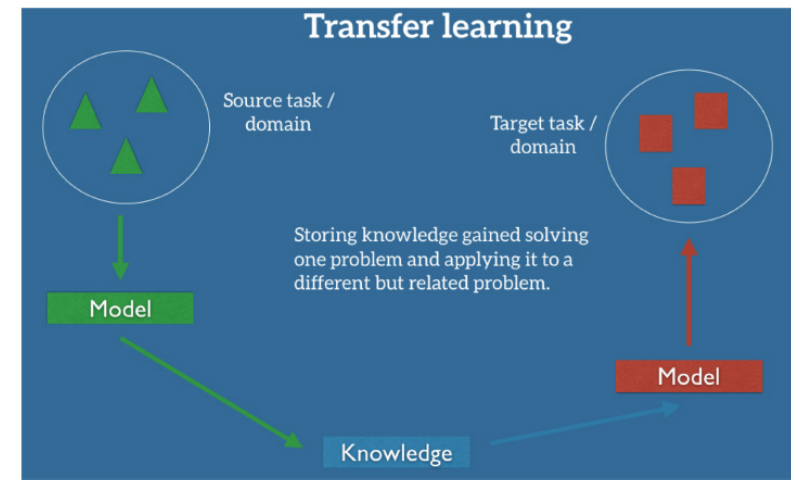
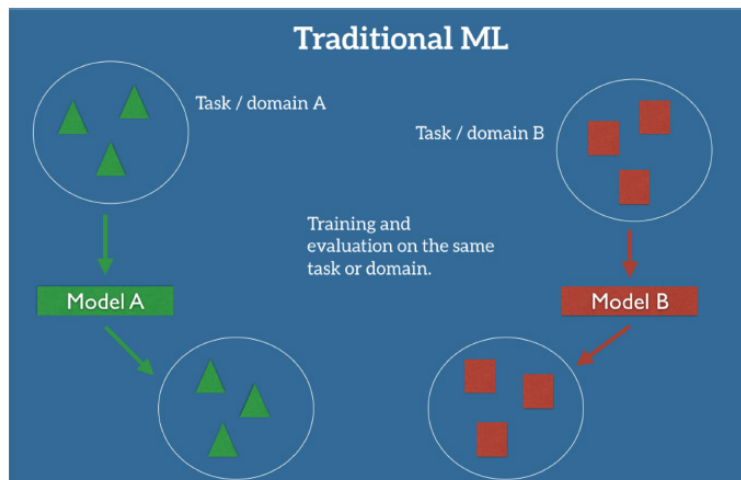
Learning Techniques

- Federated learning
 - Train on multiple decentralized devices
 - Each device has its own training data
 - Only model weights are shared and aggregated centrally
 - Important in areas where data sharing is difficult (e.g. medical information)
 - Variation: Decentralized learning, where no central aggregation is done
- Online learning
 - Update model as more observations come in time
 - Idea: Underlying data distribution changes over time & compensate for it


Learning Techniques

– Transfer learning

- Learn in one setting then exploit this to improve generalization in another setting
 - From simpler to more complex tasks
 - From a source domain (with distribution A) to another (but similar) target domain (with distribution B)

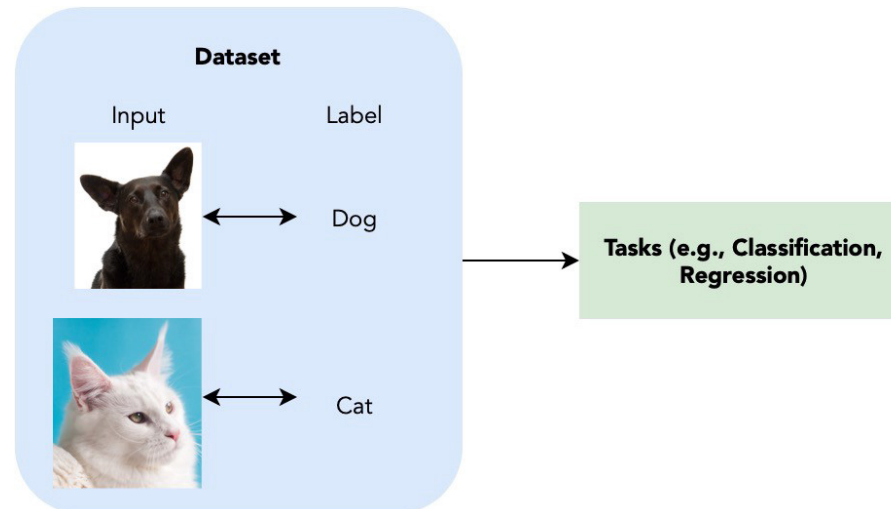


- Practically:



Learning Tasks

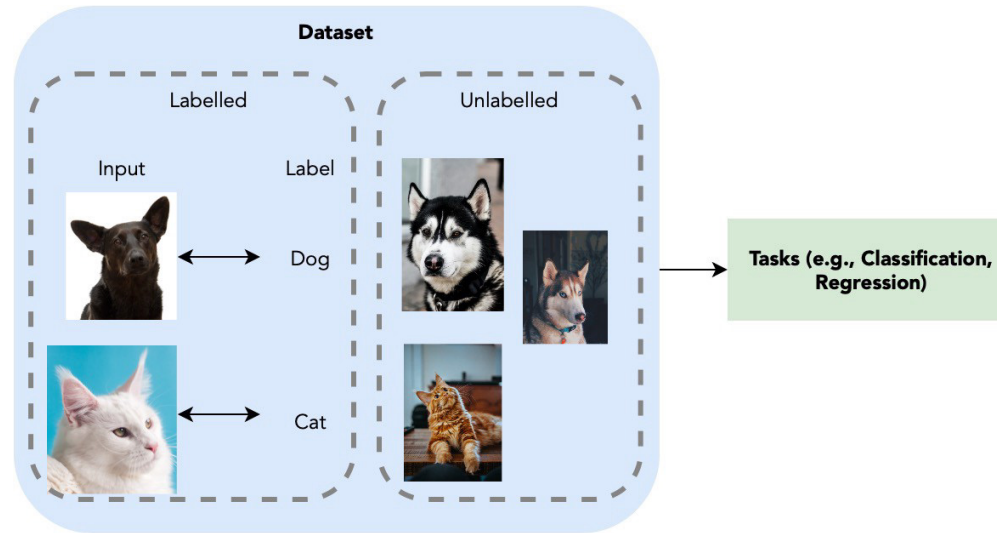
- These are problems about how we learn an algorithm based on what kind of data is available
- Supervised learning
 - Learn a mapping between input and target/label



- Issue: labels are difficult to get in real-life

Learning Tasks

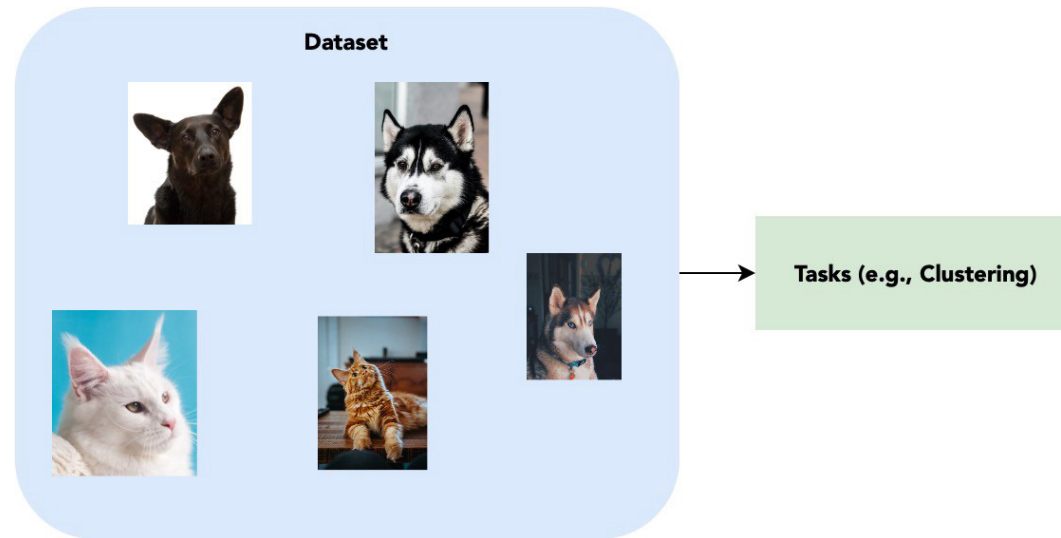
- Semi-supervised learning
 - Few labeled data + a large number of unlabeled examples



- The unlabeled data provides info about underlying data distribution, e.g.
 - Discover patterns in unlabeled data
 - Use supervised methods to label them (pseudo-labels), and assign a confidence marker to guide training

Learning Tasks

- Unsupervised learning
 - No labeled data or a-priori information

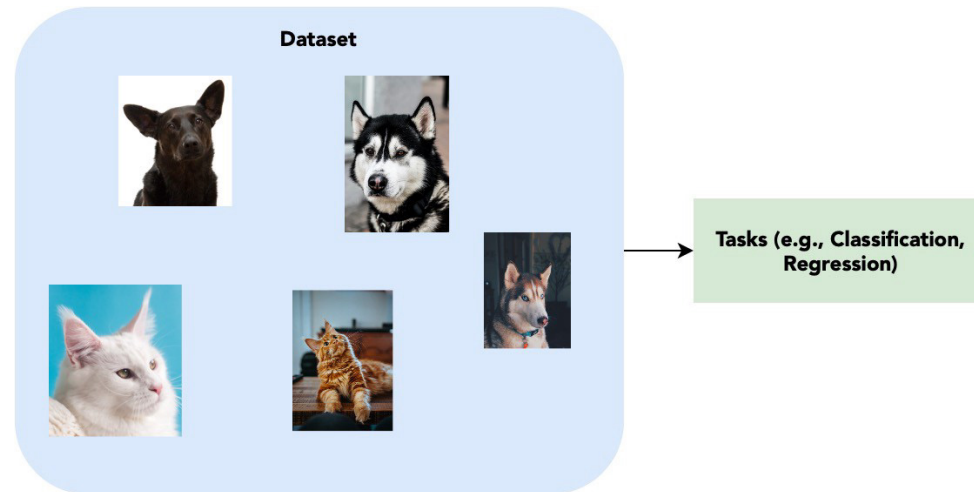


- The network learns to identify patterns in data
- Tasks are a little different, e.g. clustering, recommendation systems

Learning Tasks

- Self-supervised learning

- Circumvents problem of external labels → Still solves similar tasks to supervised learning




- Input data is used to guide the learning
- Technically unsupervised, but the loss function is framed as supervised
- Main idea: Create supervisory labels from input
 - Learn to predict some selected parts of the input from other parts of the input

Learning Tasks

– Self-supervised learning

- For instance, in regression tasks (denoising, reconstruction):
 - Mask out part of the data & try to predict it



- Same idea in natural language processing (mask out words in a sentence and predict them)
- Other ideas: Learn a different “pretext” task, fine-tune in “downstream” task (~transfer learning)
 - Here the pretext task is generated from the input data (e.g. inpainting, colorization)
 - The model is expected to learn the structure of data while learning to solve this task
 - Then apply it to downstream task
 - e.g. Contrastive learning: Do data augmentation to generate “positive” (same) and “negative” (different) input pairs → Learn to amplify features of positive pairs while hindering those of negatives

Recap

- Regularization in neural networks
 - Norm-based regularizers
 - Data augmentation
 - Early stopping
 - Dropout
- Learning Problems
 - Learning methods: Active, ensemble, federated, online, transfer learning
 - Learning tasks: Supervised, semi-supervised, unsupervised, self-supervised learning