

MNIST Classification using PyTorch


EE 5561
Fall 2023

Introduction

- Open-source machine learning framework
- Think like Numpy but with strong GPU support
- Setup:
 - `conda install pytorch torchvision -c pytorch`
 - `pip3 install torch torchvision`

MNIST Database (Modified National Institute of Standards and Technology Database)

- MNIST dataset
 - Handwritten digits
 - 60k images for training
 - 10k images for testing
 - with labels (0-9)
 - 28 × 28 pixel grayscale images

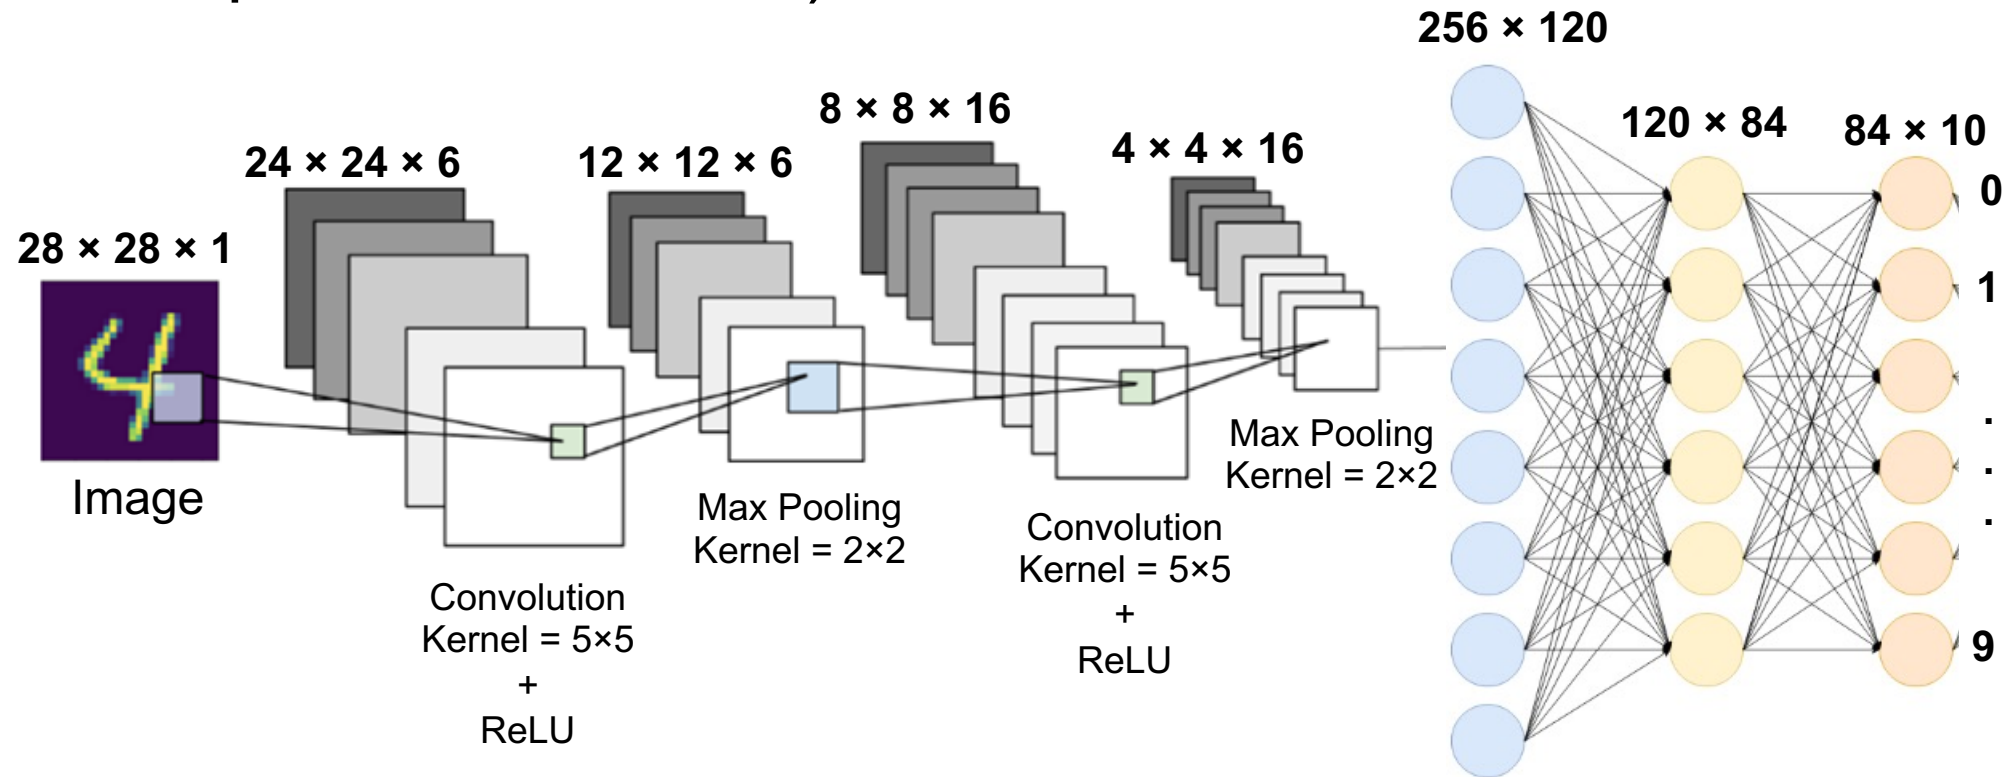


MNIST Classification

- **Task:** Classify a given handwritten image into one of 10 classes
- **Classes:** Representing the integer values from 0 to 9
- **Model:** LeNet

LeNet

- A CNN structure
- A simple well-known network used in MNIST classification (once upon a time, 1989!)



Loss Criteria

- <https://pytorch.org/docs/stable/nn.html>

TORCH.NN

These are the basic building blocks for graphs:

torch.nn

- Containers
- Convolution Layers
- Pooling layers
- Padding Layers
- Non-linear Activations (weighted sum, nonlinearity)
- Non-linear Activations (other)
- Normalization Layers
- Recurrent Layers
- Transformer Layers
- Linear Layers
- Dropout Layers
- Sparse Layers
- Distance Functions
- Loss Functions
- Vision Layers
- Shuffle Layers
- DataParallel Layers (multi-GPU, distributed)
- Utilities
- Quantized Functions
- Lazy Modules Initialization

Loss Functions

`nn.L1Loss`

Creates a criterion that measures the mean absolute error (MAE) between each element in the input x and target y .

`nn.MSELoss`

Creates a criterion that measures the mean squared error (squared L2 norm) between each element in the input x and target y .

`nn.CrossEntropyLoss`

This criterion computes the cross entropy loss between input and target.

`nn.CTCLoss`

The Connectionist Temporal Classification loss.

`nn.NLLLoss`

The negative log likelihood loss.

`nn.PoissonNLLLoss`

Negative log likelihood loss with Poisson distribution of target.

`nn.GaussianNLLLoss`

Gaussian negative log likelihood loss.

`nn.KLDivLoss`

The Kullback-Leibler divergence loss.

`nn.BCELoss`

Creates a criterion that measures the Binary Cross Entropy between the target and the input probabilities:

`nn.BCEWithLogitsLoss`

This loss combines a *Sigmoid* layer and the *BCELoss* in one

Optimizers

- <https://pytorch.org/docs/stable/optim.html>

ASGD

Implements Averaged Stochastic Gradient Descent.

LBFGS

Implements L-BFGS algorithm, heavily inspired by `minFunc`.

NAdam

Implements NAdam algorithm.

RAdam

Implements RAdam algorithm.

RMSprop

Implements RMSprop algorithm.

Rprop

Implements the resilient backpropagation algorithm.

SGD

Implements stochastic gradient descent (optionally with momentum).

Reminder

- Group project proposal deadline: Nov. 20 (**will not be extended**)
- You should be
 - Forming your groups (if you haven't done so already)
 - Coming up with project ideas
- If you do not have a group yet, email me & Merve, also include
 - What type of task you are interested in (e.g. denoising or segmentation, etc)
 - What type of application you are interested in (e.g. medical or astronomical, ...)
 - So far only heard from one person